

```

# Copyright 2018 The JAX Authors.
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
#     https://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.

```

```

import builtins
from functools import partial
import operator
from typing import Any, List, NamedTuple, Optional, Sequence, Tuple, Union

import numpy as np

```

```

from jax.interpreters import ad
from jax.interpreters import batching
from jax.interpreters import mlir
from jax._src import core
from jax._src import dtypes
from jax._src import util
from jax._src.lax import lax
from jax._src.lib import xla_client
from jax._src.lib.mlir.dialects import hlo

```

```
_max = builtins.max
```

```

Array = Any
DType = Any
Shape = core.Shape

```

```
class ConvDimensionNumbers(NamedTuple):
```

[\[docs\]](#)

```

    """Describes batch, spatial, and feature dimensions of a convolution.

    Args:
        lhs_spec: a tuple of nonnegative integer dimension numbers containing
            `(batch dimension, feature dimension, spatial dimensions...)`.
        rhs_spec: a tuple of nonnegative integer dimension numbers containing
            `(out feature dimension, in feature dimension, spatial dimensions...)`.
        out_spec: a tuple of nonnegative integer dimension numbers containing
            `(batch dimension, feature dimension, spatial dimensions...)`.
    """
    lhs_spec: Sequence[int]
    rhs_spec: Sequence[int]
    out_spec: Sequence[int]

```

```

ConvGeneralDilatedDimensionNumbers = Union[
    None, ConvDimensionNumbers, Tuple[str, str, str]]

```

```
def conv_general_dilated(
```

[\[docs\]](#)

```

    lhs: Array, rhs: Array, window_strides: Sequence[int],
    padding: Union[str, Sequence[Tuple[int, int]]],
    lhs_dilation: Optional[Sequence[int]] = None,
    rhs_dilation: Optional[Sequence[int]] = None,
    dimension_numbers: ConvGeneralDilatedDimensionNumbers = None,
    feature_group_count: int = 1, batch_group_count: int = 1,
    precision: lax.PrecisionLike = None,
    preferred_element_type: Optional[DType] = None) -> Array:
    """General n-dimensional convolution operator, with optional dilation.

```

```

    Wraps XLA's `Conv
    <https://www.tensorflow.org/xla/operation_semantics#conv_convolution>`_
    operator.

```

```
    Args:
```

```

        lhs: a rank `n+2` dimensional input array.
        rhs: a rank `n+2` dimensional array of kernel weights.
        window_strides: a sequence of `n` integers, representing the inter-window
            strides.
        padding: either the string `SAME`, the string `VALID`, or a sequence of
            `n` `(low, high)` integer pairs that give the padding to apply before and
            after each spatial dimension.
        lhs_dilation: `None`, or a sequence of `n` integers, giving the
            dilation factor to apply in each spatial dimension of `lhs`. LHS dilation
            is also known as transposed convolution.
        rhs_dilation: `None`, or a sequence of `n` integers, giving the
            dilation factor to apply in each spatial dimension of `rhs`. RHS dilation
            is also known as atrous convolution.
        dimension_numbers: either `None`, a `ConvDimensionNumbers` object, or
            a 3-tuple `(lhs_spec, rhs_spec, out_spec)`, where each element is a
            string of length `n+2`.

```

feature_group_count: integer, default 1. See XLA HLO docs.
 batch_group_count: integer, default 1. See XLA HLO docs.
 precision: Optional. Either ``None``, which means the default precision for the backend, a :class:`~jax.lax.Precision` enum value (``Precision.DEFAULT``, ``Precision.HIGH`` or ``Precision.HIGHEST``), a string (e.g. 'highest' or 'fastest', see the ``jax.default_matmul_precision`` context manager), or a tuple of two :class:`~jax.lax.Precision` enums or strings indicating precision of
 ``lhs`` and ``rhs``.
 preferred_element_type: Optional. Either ``None``, which means the default accumulation type for the input types, or a datatype, indicating to accumulate results to and return a result with that datatype.

Returns:

An array containing the convolution result.

In the string case of ``dimension_numbers``, each character identifies by position:

- the batch dimensions in ``lhs``, ``rhs``, and the output with the character 'N',
- the feature dimensions in ``lhs`` and the output with the character 'C',
- the input and output feature dimensions in ``rhs`` with the characters 'I' and 'O' respectively, and
- spatial dimension correspondences between ``lhs``, ``rhs``, and the output using any distinct characters.

For example, to indicate dimension numbers consistent with the ``conv`` function with two spatial dimensions, one could use ``('NCHW', 'OIHW', 'NCHW')``. As another example, to indicate dimension numbers consistent with the TensorFlow Conv2D operation, one could use ``('NHWC', 'HWIO', 'NHWC')``. When using the latter form of convolution dimension specification, window strides are associated with spatial dimension character labels according to the order in which the labels appear in the ``rhs_spec`` string, so that ``window_strides[0]`` is matched with the dimension corresponding to the first character appearing in ``rhs_spec`` that is not ``'I'`` or ``'O'``.

If ``dimension_numbers`` is ``None``, the default is ``('NCHW', 'OIHW', 'NCHW')`` (for a 2D convolution).

```

"""
dnums = conv_dimension_numbers(lhs.shape, rhs.shape, dimension_numbers)
if lhs_dilation is None:
    lhs_dilation = (1,) * (lhs.ndim - 2)
elif isinstance(padding, str) and not len(lhs_dilation) == lhs_dilation.count(1):
    raise ValueError(
        "String padding is not implemented for transposed convolution "
        "using this op. Please either exactly specify the required padding or "
        "use conv_transpose.")
if rhs_dilation is None:
    rhs_dilation = (1,) * (rhs.ndim - 2)
if isinstance(padding, str):
    lhs_perm, rhs_perm, _ = dnums
    rhs_shape = np.take(rhs.shape, rhs_perm)[2:] # type: ignore[index]
    effective_rhs_shape = [(k-1) * r + 1 for k, r in zip(rhs_shape, rhs_dilation)]
    padding = lax.padtype_to_pads(
        np.take(lhs.shape, lhs_perm)[2:], effective_rhs_shape, # type:
ignore[index]
        window_strides, padding)
else:
    try:
        padding = tuple((operator.index(lo), operator.index(hi))
                        for lo, hi in padding)
    except (ValueError, TypeError) as e:
        raise ValueError(
            "padding to conv_general_dilated should be a string or a "
            f"sequence of (low, high) pairs, got {padding}") from e

preferred_element_type = (
    None if preferred_element_type is None else
    dtypes.canonicalize_dtype(np.dtype(preferred_element_type)))
return conv_general_dilated_p.bind(
    lhs, rhs, window_strides=tuple(window_strides), padding=tuple(padding),
    lhs_dilation=tuple(lhs_dilation), rhs_dilation=tuple(rhs_dilation),
    dimension_numbers=dnums,
    feature_group_count=feature_group_count,
    batch_group_count=batch_group_count,
    lhs_shape=lhs.shape, rhs_shape=rhs.shape,
    precision=lax.canonicalize_precision(precision),
    preferred_element_type=preferred_element_type)

```

convenience wrappers around traceables

```

def conv(lhs: Array, rhs: Array, window_strides: Sequence[int],
        padding: str, precision: lax.PrecisionLike = None,
        preferred_element_type: Optional[DType] = None) -> Array:
    """Convenience wrapper around `conv_general_dilated`.

```

[\[docs\]](#)

```

Args:
    lhs: a rank `n+2` dimensional input array.
    rhs: a rank `n+2` dimensional array of kernel weights.
    window_strides: a sequence of `n` integers, representing the inter-window
        strides.
    padding: either the string `'SAME'`, the string `'VALID'`.
    precision: Optional. Either ``None``, which means the default precision for
        the backend, a :class:`~jax.lax.Precision` enum value (``Precision.DEFAULT``,
        ``Precision.HIGH`` or ``Precision.HIGHEST``) or a tuple of two
        :class:`~jax.lax.Precision` enums indicating precision of ``lhs`` and
        ``rhs``.
    preferred_element_type: Optional. Either ``None``, which means the default
        accumulation type for the input types, or a datatype, indicating to
        accumulate results to and return a result with that datatype.

Returns:
    An array containing the convolution result.
    """
    return conv_general_dilated(lhs, rhs, window_strides, padding,
                                precision=precision,
                                preferred_element_type=preferred_element_type)

def conv_with_general_padding(lhs: Array, rhs: Array,                                [docs]
                              window_strides: Sequence[int],
                              padding: Union[str, Sequence[Tuple[int, int]]],
                              lhs_dilation: Optional[Sequence[int]],
                              rhs_dilation: Optional[Sequence[int]],
                              precision: lax.PrecisionLike = None,
                              preferred_element_type: Optional[DType] = None) ->
Array:
    """Convenience wrapper around `conv_general_dilated`.

Args:
    lhs: a rank `n+2` dimensional input array.
    rhs: a rank `n+2` dimensional array of kernel weights.
    window_strides: a sequence of `n` integers, representing the inter-window
        strides.
    padding: either the string `'SAME'`, the string `'VALID'`, or a sequence of
        `n` `(low, high)` integer pairs that give the padding to apply before and
        after each spatial dimension.
    lhs_dilation: `None`, or a sequence of `n` integers, giving the
        dilation factor to apply in each spatial dimension of `lhs`. LHS dilation
        is also known as transposed convolution.
    rhs_dilation: `None`, or a sequence of `n` integers, giving the
        dilation factor to apply in each spatial dimension of `rhs`. RHS dilation
        is also known as atrous convolution.
    precision: Optional. Either ``None``, which means the default precision for
        the backend, a :class:`~jax.lax.Precision` enum value (``Precision.DEFAULT``,
        ``Precision.HIGH`` or ``Precision.HIGHEST``) or a tuple of two
        :class:`~jax.lax.Precision` enums indicating precision of ``lhs`` and
        ``rhs``.
    preferred_element_type: Optional. Either ``None``, which means the default
        accumulation type for the input types, or a datatype, indicating to
        accumulate results to and return a result with that datatype.

Returns:
    An array containing the convolution result.
    """
    return conv_general_dilated(
        lhs, rhs, window_strides, padding, lhs_dilation=lhs_dilation,
        rhs_dilation=rhs_dilation, precision=precision,
        preferred_element_type=preferred_element_type)

def _conv_transpose_padding(k, s, padding):
    """Calculate before and after padding for a dim of transposed convolution.

Args:
    k: int: kernel dimension.
    s: int: dimension stride value.
    padding: 'same' or 'valid' padding mode for original forward conv.

Returns:
    2-tuple: ints: before and after padding for transposed convolution.
    """
    if padding == 'SAME':
        pad_len = k + s - 2
        if s > k - 1:
            pad_a = k - 1
        else:
            pad_a = int(np.ceil(pad_len / 2))
    elif padding == 'VALID':
        pad_len = k + s - 2 + _max(k - s, 0)
        pad_a = k - 1
    else:
        raise ValueError('Padding mode must be `SAME` or `VALID`.')
    pad_b = pad_len - pad_a
    return pad_a, pad_b

```

```

def _flip_axes(x, axes):
    """Flip ndarray 'x' along each axis specified in axes tuple."""
    for axis in axes:
        x = np.flip(x, axis)
    return x

def conv_transpose(lhs: Array, rhs: Array, strides: Sequence[int], padding: Union[str, Sequence[Tuple[int, int]]], rhs_dilation: Optional[Sequence[int]] = None, dimension_numbers: ConvGeneralDilatedDimensionNumbers = None, transpose_kernel: bool = False, precision: lax.PrecisionLike = None, preferred_element_type: Optional[DType] = None) -> Array:
    """Convenience wrapper for calculating the N-d convolution "transpose".

    This function directly calculates a fractionally strided conv rather than indirectly calculating the gradient (transpose) of a forward convolution.

    Args:
        lhs: a rank 'n+2' dimensional input array.
        rhs: a rank 'n+2' dimensional array of kernel weights.
        strides: sequence of 'n' integers, sets fractional stride.
        padding: 'SAME', 'VALID' will set as transpose of corresponding forward conv, or a sequence of 'n' integer 2-tuples describing before-and-after padding for each 'n' spatial dimension.
        rhs_dilation: 'None', or a sequence of 'n' integers, giving the dilation factor to apply in each spatial dimension of 'rhs'. RHS dilation is also known as atrous convolution.
        dimension_numbers: tuple of dimension descriptors as in lax.conv_general_dilated. Defaults to tensorflow convention.
        transpose_kernel: if True flips spatial axes and swaps the input/output channel axes of the kernel. This makes the output of this function identical to the gradient-derived functions like keras.layers.Conv2DTranspose applied to the same kernel. For typical use in neural nets this is completely pointless and just makes input/output channel specification confusing.
        precision: Optional. Either ``None``, which means the default precision for the backend, a :class:`~jax.lax.Precision` enum value (``Precision.DEFAULT``, ``Precision.HIGH`` or ``Precision.HIGHEST``) or a tuple of two :class:`~jax.lax.Precision` enums indicating precision of ``lhs`` and ``rhs``.
        preferred_element_type: Optional. Either ``None``, which means the default accumulation type for the input types, or a datatype, indicating to accumulate results to and return a result with that datatype.

    Returns:
        Transposed N-d convolution, with output padding following the conventions of keras.layers.Conv2DTranspose.
    """
    assert len(lhs.shape) == len(rhs.shape) and len(lhs.shape) >= 2
    ndims = len(lhs.shape)
    one = (1,) * (ndims - 2)
    # Set dimensional layout defaults if not specified.
    if dimension_numbers is None:
        if ndims == 2:
            dimension_numbers = ('NC', 'IO', 'NC')
        elif ndims == 3:
            dimension_numbers = ('NHC', 'HIO', 'NHC')
        elif ndims == 4:
            dimension_numbers = ('NHWC', 'HWIO', 'NHWC')
        elif ndims == 5:
            dimension_numbers = ('NHWDC', 'HWDIO', 'NHWDC')
        else:
            raise ValueError('No 4+ dimensional dimension_number defaults.')
    dn = conv_dimension_numbers(lhs.shape, rhs.shape, dimension_numbers)
    k_shape = np.take(rhs.shape, dn.rhs_spec)
    k_sdims = k_shape[2:] # type: ignore[index]
    # Calculate correct output shape given padding and strides.
    pads: Union[str, Sequence[Tuple[int, int]]]
    if isinstance(padding, str) and padding in {'SAME', 'VALID'}:
        if rhs_dilation is None:
            rhs_dilation = (1,) * (rhs.ndim - 2)
            effective_k_size = map(lambda k, r: (k-1) * r + 1, k_sdims, rhs_dilation)
            pads = [_conv_transpose_padding(k, s, padding) for k,s in zip(effective_k_size, strides)]
        else:
            pads = padding
    if transpose_kernel:
        # flip spatial dims and swap input / output channel axes
        rhs = _flip_axes(rhs, np.array(dn.rhs_spec)[2:])
        rhs = np.swapaxes(rhs, dn.rhs_spec[0], dn.rhs_spec[1])
    return conv_general_dilated(lhs, rhs, one, pads, strides, rhs_dilation, dn, precision=precision, preferred_element_type=preferred_element_type)

def _conv_general_dilated_shape_rule(

```

```

lhs: core.ShapedArray, rhs: core.ShapedArray, *, window_strides, padding,
lhs_dilation, rhs_dilation, dimension_numbers, feature_group_count,
batch_group_count, **unused_kwargs) -> Tuple[int, ...]:
assert type(dimension_numbers) is ConvDimensionNumbers
if len(lhs.shape) != len(rhs.shape):
    msg = ("conv_general_dilated lhs and rhs must have the same number of "
           "dimensions, but got {} and {}.")
    raise ValueError(msg.format(lhs.shape, rhs.shape))
if not feature_group_count > 0:
    msg = ("conv_general_dilated feature_group_count "
           "must be a positive integer, got {}.")
    raise ValueError(msg.format(feature_group_count))
lhs_feature_count = lhs.shape[dimension_numbers.lhs_spec[1]]
quot, rem = divmod(lhs_feature_count, feature_group_count)
if rem:
    msg = ("conv_general_dilated feature_group_count must divide lhs feature "
           "dimension size, but {} does not divide {}.")
    raise ValueError(msg.format(feature_group_count, lhs_feature_count))
if not core.symbolic_equal_dim(quot, rhs.shape[dimension_numbers.rhs_spec[1]]):
    msg = ("conv_general_dilated lhs feature dimension size divided by "
           "feature_group_count must equal the rhs input feature dimension "
           "size, but {} // {} != {}.")
    raise ValueError(msg.format(lhs_feature_count, feature_group_count,
                                rhs.shape[dimension_numbers.rhs_spec[1]]))
if rhs.shape[dimension_numbers.rhs_spec[0]] % feature_group_count:
    msg = ("conv_general_dilated rhs output feature dimension size must be a "
           "multiple of feature_group_count, but {} is not a multiple of {}.")
    raise ValueError(msg.format(rhs.shape[dimension_numbers.rhs_spec[0]],
                                feature_group_count))

if not batch_group_count > 0:
    msg = ("conv_general_dilated batch_group_count "
           "must be a positive integer, got {}.")
    raise ValueError(msg.format(batch_group_count))
lhs_batch_count = lhs.shape[dimension_numbers.lhs_spec[0]]
if batch_group_count > 1 and lhs_batch_count % batch_group_count != 0:
    msg = ("conv_general_dilated batch_group_count must divide lhs batch "
           "dimension size, but {} does not divide {}.")
    raise ValueError(msg.format(batch_group_count, lhs_batch_count))

if rhs.shape[dimension_numbers.rhs_spec[0]] % batch_group_count:
    msg = ("conv_general_dilated rhs output feature dimension size must be a "
           "multiple of batch_group_count, but {} is not a multiple of {}.")
    raise ValueError(msg.format(rhs.shape[dimension_numbers.rhs_spec[0]],
                                batch_group_count))

if batch_group_count > 1 and feature_group_count > 1:
    msg = ("At most one of batch_group_count and feature_group_count may be > "
           "1, got batch_group_count={} and feature_group_count={}")
    raise ValueError(msg.format(batch_group_count, feature_group_count))

if len(_conv_sdims(dimension_numbers.rhs_spec)) != len(window_strides):
    msg = ("conv_general_dilated window and window_strides must have "
           "the same number of dimensions, but got {} and {}.")
    raise ValueError(
        msg.format(len(_conv_sdims(dimension_numbers.rhs_spec)),
                    len(window_strides)))

lhs_perm, rhs_perm, out_perm = dimension_numbers
lhs_trans = lax.dilate_shape(np.take(lhs.shape, lhs_perm), lhs_dilation)
rhs_trans = lax.dilate_shape(np.take(rhs.shape, rhs_perm), rhs_dilation)
out_trans = conv_shape_tuple(lhs_trans, rhs_trans, window_strides, padding,
                              batch_group_count)
return tuple(np.take(out_trans, np.argsort(out_perm))) # type: ignore[arg-type]

def _conv_general_dilated_dtype_rule(
    lhs, rhs, *, window_strides, padding, lhs_dilation, rhs_dilation,
    dimension_numbers, preferred_element_type, **unused_kwargs):
    input_dtype = lax.naryop_dtype_rule(lax._input_dtype, [lax._any, lax._any],
                                         'conv_general_dilated', lhs, rhs)

    if preferred_element_type is None:
        return input_dtype
    lax._validate_preferred_element_type(input_dtype, preferred_element_type)
    return preferred_element_type

_conv_spec_transpose = lambda spec: (spec[1], spec[0]) + spec[2:]
_conv_sdims = lambda spec: spec[2:]

```

```

# Understanding the convolution transpose rules:
# Ignoring the spatial dimensions, let m = batch, j = input feature,
# k = output feature.
#
# Convolution computes the following contraction:
# Forward: [m, j] [j, k] -> [m, k]
#
# The transposes are similar to the rules for transposing a matmul:
# LHS transpose: [m, k] [k, j] -> [m, j]
# RHS transpose: [j, m] [m, k] -> [j, k]
#

```

```

# With feature grouping, we have the following signatures:
# Forward: [m, gj] [j, gk] -> [m, gk]
# LHS transpose: [m, gk] [k, gj] -> [m, gj]
# --> implemented as feature grouping after transposing the group from the
#     kernel input features to the kernel output features.
# RHS transpose: [gj, m] [m, gk] -> [j, gk]
# --> which is batch grouping.
#
# With batch grouping, we have the following signatures:
# Forward: [gm, j] [j, gk] -> [m, gk]
# LHS transpose: [m, gk] [gk, j] -> [gm, j]
# --> implemented as feature grouping with transposing the group on the kernel
#     and the output.
# RHS transpose: [j, gm] [m, gk] -> [j, gk]
# --> which is feature grouping.

```

```

def _conv_general_dilated_transpose_lhs(
    g, rhs, *, window_strides, padding, lhs_dilation, rhs_dilation,
    dimension_numbers, feature_group_count, batch_group_count,
    lhs_shape, rhs_shape, precision, preferred_element_type):
    assert type(dimension_numbers) is ConvDimensionNumbers
    assert batch_group_count == 1 or feature_group_count == 1
    lhs_sdims, rhs_sdims, out_sdims = map(_conv_sdims, dimension_numbers)
    lhs_spec, rhs_spec, out_spec = dimension_numbers
    t_rhs_spec = _conv_spec_transpose(rhs_spec)
    if feature_group_count > 1:
        # in addition to switching the dims in the spec, need to move the feature
        # group axis into the transposed rhs's output feature dim
        rhs = _reshape_axis_out_of(rhs_spec[0], feature_group_count, rhs)
        rhs = _reshape_axis_into(rhs_spec[0], rhs_spec[1], rhs)
    elif batch_group_count > 1:
        rhs = _reshape_axis_out_of(rhs_spec[0], batch_group_count, rhs)
        rhs = _reshape_axis_into(rhs_spec[0], rhs_spec[1], rhs)
        feature_group_count = batch_group_count
    trans_dimension_numbers = ConvDimensionNumbers(out_spec, t_rhs_spec, lhs_spec)
    padding = _conv_general_vjp_lhs_padding(
        np.take(lhs_shape, lhs_sdims), np.take(rhs_shape, rhs_sdims),
        window_strides, np.take(g.shape, out_sdims), padding, lhs_dilation,
        rhs_dilation)
    revd_weights = lax.rev(rhs, rhs_sdims)
    out = conv_general_dilated(
        g, revd_weights, window_strides=lhs_dilation, padding=padding,
        lhs_dilation=window_strides, rhs_dilation=rhs_dilation,
        dimension_numbers=trans_dimension_numbers,
        feature_group_count=feature_group_count,
        batch_group_count=1, precision=precision,
        preferred_element_type=preferred_element_type)
    if batch_group_count > 1:
        out = _reshape_axis_out_of(lhs_spec[1], batch_group_count, out)
        out = _reshape_axis_into(lhs_spec[1], lhs_spec[0], out)
    return out

def _conv_general_dilated_transpose_rhs(
    g, lhs, *, window_strides, padding, lhs_dilation, rhs_dilation,
    dimension_numbers: ConvDimensionNumbers, feature_group_count: int,
    batch_group_count: int, lhs_shape, rhs_shape, precision,
    preferred_element_type):
    assert type(dimension_numbers) is ConvDimensionNumbers
    if np.size(g) == 0:
        # Avoids forming degenerate convolutions where the RHS has spatial size 0.
        # Awkwardly, we don't have an aval for the rhs readily available, so instead
        # of returning an ad_util.Zero instance here, representing a symbolic zero
        # value, we instead return a None, which is meant to represent having no
        # cotangent at all (and is thus incorrect for this situation), since the two
        # are treated the same operationally.
        # TODO(mattjj): adjust defbilinear so that the rhs aval is available here
        return None
    lhs_sdims, rhs_sdims, out_sdims = map(_conv_sdims, dimension_numbers)
    lhs_trans, rhs_trans, out_trans = map(_conv_spec_transpose, dimension_numbers)
    assert batch_group_count == 1 or feature_group_count == 1
    if batch_group_count > 1:
        feature_group_count = batch_group_count
        batch_group_count = 1
    elif feature_group_count > 1:
        batch_group_count = feature_group_count
        feature_group_count = 1
    trans_dimension_numbers = ConvDimensionNumbers(lhs_trans, out_trans, rhs_trans)
    padding = _conv_general_vjp_rhs_padding(
        np.take(lhs_shape, lhs_sdims), np.take(rhs_shape, rhs_sdims),
        window_strides, np.take(g.shape, out_sdims), padding, lhs_dilation,
        rhs_dilation)
    return conv_general_dilated(
        lhs, g, window_strides=rhs_dilation, padding=padding,
        lhs_dilation=lhs_dilation, rhs_dilation=window_strides,
        dimension_numbers=trans_dimension_numbers,
        feature_group_count=feature_group_count,
        batch_group_count=batch_group_count, precision=precision,
        preferred_element_type=preferred_element_type)

def _conv_general_dilated_batch_rule(

```



```

batched_args, batch_dims, *, window_strides, padding,
lhs_dilation, rhs_dilation, dimension_numbers,
feature_group_count, batch_group_count, precision,
preferred_element_type, **unused_kwargs):
assert batch_group_count == 1 or feature_group_count == 1
lhs, rhs = batched_args
lhs_bdim, rhs_bdim = batch_dims
lhs_spec, rhs_spec, out_spec = dimension_numbers

# Some of the cases that reshape into batch or feature dimensions do not work
# with size 0 batch dimensions. The best fix would be to extend HLO to support
# multiple batch dimensions.
if ((lhs_bdim is not None and lhs.shape[lhs_bdim] == 0) or
    (rhs_bdim is not None and rhs.shape[rhs_bdim] == 0)):
    lhs_shape_unbatched, rhs_shape_unbatched = list(lhs.shape), list(rhs.shape)
    if lhs_bdim is not None:
        lhs_shape_unbatched.pop(lhs_bdim)
    if rhs_bdim is not None:
        rhs_shape_unbatched.pop(rhs_bdim)
    shape = _conv_general_dilated_shape_rule(
        core.ShapedArray(lhs_shape_unbatched, lhs.dtype),
        core.ShapedArray(rhs_shape_unbatched, rhs.dtype),
        window_strides=window_strides, padding=padding, lhs_dilation=lhs_dilation,
        rhs_dilation=rhs_dilation, dimension_numbers=dimension_numbers,
        feature_group_count=feature_group_count,
        batch_group_count=batch_group_count)
    return lax.full(
        (0,) + shape, 0,
        dtype=lhs.dtype if preferred_element_type is None
            else preferred_element_type), 0

if lhs_bdim is not None and rhs_bdim is not None:
    assert lhs.shape[lhs_bdim] == rhs.shape[rhs_bdim]
    if batch_group_count > 1:
        new_lhs = _reshape_axis_into(lhs_bdim, lhs_spec[0], lhs)
        batch_group_count *= lhs.shape[lhs_bdim]
    else:
        new_lhs = _reshape_axis_into(lhs_bdim, lhs_spec[1], lhs)
        feature_group_count *= lhs.shape[lhs_bdim]
    new_rhs = _reshape_axis_into(rhs_bdim, rhs_spec[0], rhs)
    out = conv_general_dilated(
        new_lhs, new_rhs, window_strides, padding, lhs_dilation, rhs_dilation,
        dimension_numbers, feature_group_count=feature_group_count,
        batch_group_count=batch_group_count, precision=precision,
        preferred_element_type=preferred_element_type)
    out = _reshape_axis_out_of(out_spec[1], lhs.shape[lhs_bdim], out)
    return out, out_spec[1]

elif lhs_bdim is not None:
    if batch_group_count == 1:
        new_lhs = _reshape_axis_into(lhs_bdim, lhs_spec[0], lhs)
        out = conv_general_dilated(new_lhs, rhs, window_strides, padding,
                                   lhs_dilation, rhs_dilation, dimension_numbers,
                                   feature_group_count, precision=precision,
                                   preferred_element_type=preferred_element_type)
        out = _reshape_axis_out_of(out_spec[0], lhs.shape[lhs_bdim], out)
        return out, out_spec[0]
    else:
        new_lhs = _reshape_axis_out_of(lhs_spec[0] + int(lhs_bdim <= lhs_spec[0]),
                                       batch_group_count, lhs)
        new_lhs = _reshape_axis_into(lhs_bdim + int(lhs_spec[0] < lhs_bdim),
                                     lhs_spec[0] + 1,
                                     new_lhs)
        new_lhs = _reshape_axis_into(lhs_spec[0], lhs_spec[0], new_lhs)
        out = conv_general_dilated(new_lhs, rhs, window_strides, padding,
                                   lhs_dilation, rhs_dilation, dimension_numbers,
                                   feature_group_count, batch_group_count,
                                   precision=precision,
                                   preferred_element_type=preferred_element_type)
        out = _reshape_axis_out_of(out_spec[0], lhs.shape[lhs_bdim], out)
        return out, out_spec[0]

elif rhs_bdim is not None:
    if feature_group_count == 1 and batch_group_count == 1:
        new_rhs = _reshape_axis_into(rhs_bdim, rhs_spec[0], rhs)
        out = conv_general_dilated(lhs, new_rhs, window_strides, padding,
                                   lhs_dilation, rhs_dilation, dimension_numbers,
                                   feature_group_count, batch_group_count,
                                   precision=precision,
                                   preferred_element_type=preferred_element_type)
        out = _reshape_axis_out_of(out_spec[1], rhs.shape[rhs_bdim], out)
        return out, out_spec[1]
    else:
        # groups need to be outermost, so we need to factor them out of the
        # rhs output feature dim, then factor the batch dim into the remaining rhs
        # output feature dim, then put groups back in. We do something
        # similar on the output. An alternative which would require more FLOPs but
        # fewer reshapes would be to broadcast lhs.
        group_count = (feature_group_count if feature_group_count > 1

```

```

        else batch_group_count)
    new_rhs = _reshape_axis_out_of(rhs_spec[0] + int(rhs_bdim <= rhs_spec[0]),
                                   group_count, rhs)
    new_rhs = _reshape_axis_into(rhs_bdim + int(rhs_spec[0] < rhs_bdim),
                                 rhs_spec[0] + 1, new_rhs)
    new_rhs = _reshape_axis_into(rhs_spec[0], rhs_spec[0], new_rhs)
    out = conv_general_dilated(lhs, new_rhs, window_strides, padding,
                               lhs_dilation, rhs_dilation, dimension_numbers,
                               feature_group_count, batch_group_count,
                               precision=precision,
                               preferred_element_type=preferred_element_type)
    out = _reshape_axis_out_of(out_spec[1], group_count, out)
    out = _reshape_axis_out_of(out_spec[1] + 1, rhs.shape[rhs_bdim], out)
    out = _reshape_axis_into(out_spec[1], out_spec[1] + 1, out)
    return out, out_spec[1]

conv_general_dilated_p = lax.standard_primitive(
    _conv_general_dilated_shape_rule, _conv_general_dilated_dtype_rule,
    'conv_general_dilated')

ad.defbilinear(conv_general_dilated_p,
               _conv_general_dilated_transpose_lhs,
               _conv_general_dilated_transpose_rhs)
batching.primitive_batchers[conv_general_dilated_p] = \
    _conv_general_dilated_batch_rule

def _complex_mul(mul, x, y):
    # We use a trick for complex multiplication sometimes attributed to Gauss
    # which uses three multiplications and five additions; instead of the naive
    # method of four multiplications and two additions.
    #
    https://en.wikipedia.org/wiki/Multiplication\_algorithm#Complex\_multiplication\_algorithm
    #
    # This performance win comes with a trade-off in accuracy; especially in
    # cases when the real and imaginary differ hugely in magnitude. The relative
    # error bound (e.g. 1p-24 in case of float32) would be relative to the
    # maximum of real and imaginary parts of the result instead of being
    # satisfied by the real and imaginary parts independently of each other.
    x_re, x_im = lax.real(x), lax.imag(x)
    y_re, y_im = lax.real(y), lax.imag(y)
    k1 = mul(lax.add(x_re, x_im), y_re)
    k2 = mul(x_re, lax.sub(y_im, y_re))
    k3 = mul(x_im, lax.add(y_re, y_im))
    return lax.complex(lax.sub(k1, k3), lax.add(k1, k2))

_real_dtype = lambda dtype: np.finfo(dtype).dtype

def _conv_general_dilated_lower(
    ctx, lhs, rhs, *, window_strides, padding,
    lhs_dilation, rhs_dilation, dimension_numbers, feature_group_count,
    batch_group_count, precision, preferred_element_type,
    expand_complex_convolutions=False, **unused_kwargs):
    lhs_aval, rhs_aval = ctx.aval_in
    aval_out, = ctx.aval_out
    assert isinstance(dimension_numbers, ConvDimensionNumbers)
    dtype = lhs_aval.dtype
    if expand_complex_convolutions and np.issubdtype(dtype, np.complexfloating):
        if preferred_element_type is not None:
            # Convert complex dtype to types used for real and imaginary parts
            assert np.issubdtype(preferred_element_type, np.complexfloating)
            preferred_element_type = _real_dtype(preferred_element_type)
        complex_conv = mlir.lower_fun(
            partial(
                _complex_mul,
                partial(conv_general_dilated, window_strides=window_strides,
                    padding=padding, lhs_dilation=lhs_dilation,
                    rhs_dilation=rhs_dilation, dimension_numbers=dimension_numbers,
                    feature_group_count=feature_group_count,
                    batch_group_count=batch_group_count, precision=precision,
                    preferred_element_type=preferred_element_type))),
            multiple_results=False)
        return complex_conv(ctx, lhs, rhs)

lhs_spec, rhs_spec, out_spec = dimension_numbers
dnums = hlo.ConvDimensionNumbers.get(
    input_batch_dimension=lhs_spec[0],
    input_feature_dimension=lhs_spec[1],
    input_spatial_dimensions=list(lhs_spec[2:]),
    kernel_output_feature_dimension=rhs_spec[0],
    kernel_input_feature_dimension=rhs_spec[1],
    kernel_spatial_dimensions=list(rhs_spec[2:]),
    output_batch_dimension=out_spec[0],
    output_feature_dimension=out_spec[1],
    output_spatial_dimensions=list(out_spec[2:]))
num_spatial_dims = len(rhs_spec) - 2
if len(padding) == 0:
    padding = np.zeros((0, 2), dtype=np.int64)
window_reversal = mlir.dense_bool_elements([False] * num_spatial_dims)

```



```

    return [
        hlo.ConvolutionOp(
            mlir.aval_to_ir_type(aval_out),
            lhs,
            rhs,
            dimension_numbers=dnums,
            feature_group_count=mlir.i64_attr(feature_group_count),
            batch_group_count=mlir.i64_attr(batch_group_count),
            window_strides=mlir.dense_int_elements(window_strides),
            padding=mlir.dense_int_elements(padding),
            lhs_dilation=mlir.dense_int_elements(lhs_dilation),
            rhs_dilation=mlir.dense_int_elements(rhs_dilation),
            window_reversal=window_reversal,
            precision_config=lax.precision_attr(precision)).result
    ]

mlir.register_lowering(conv_general_dilated_p, _conv_general_dilated_lower)
# TODO(b/161124619, b/161126248): XLA does not support complex convolution on
# GPU, and on CPU it uses a slow loop-based implementation;
# on these backends, lower complex convolutions away.
mlir.register_lowering(
    conv_general_dilated_p,
    partial(_conv_general_dilated_lower, expand_complex_convolutions=True),
    platform='cpu')
mlir.register_lowering(
    conv_general_dilated_p,
    partial(_conv_general_dilated_lower, expand_complex_convolutions=True),
    platform='gpu')

def _reshape_axis_into(src, dst, x):
    # NB: `dst` is the number of the dimension that we should reshape into
    # *after* `src` is removed from `x`'s list of dimensions. For example, if
    # `src` is an added batch dimension, `dst` might name a target dimension in
    # the unbatched list of dimensions.
    perm = [i for i in range(x.ndim) if i != src]
    perm.insert(dst, src)
    new_shape = list(np.delete(x.shape, src))
    new_shape[dst] *= x.shape[src]
    return lax.reshape(x, new_shape, perm)

def _reshape_axis_out_of(src, size1, x):
    shape = list(x.shape)
    size2, ragged = divmod(shape[src], size1)
    assert not ragged
    shape[src:src+1] = [size1, size2]
    return lax.reshape(x, shape)

def _check_conv_shapes(name, lhs_shape, rhs_shape, window_strides):
    """Check that conv shapes are valid and are consistent with window_strides."""
    if len(lhs_shape) != len(rhs_shape):
        msg = "Arguments to {} must have same rank, got {} and {}."
        raise TypeError(msg.format(name, len(lhs_shape), len(rhs_shape)))
    if len(lhs_shape) < 2:
        msg = "Arguments to {} must have rank at least 2, got {} and {}."
        raise TypeError(msg.format(name, len(lhs_shape), len(rhs_shape)))
    if lhs_shape[1] != rhs_shape[1]:
        msg = "Arguments to {} must agree on input feature size, got {} and {}."
        raise TypeError(msg.format(name, lhs_shape[1], rhs_shape[1]))
    lax._check_shape_like(name, "window_strides", window_strides)
    if not np.all(np.greater(window_strides, 0)):
        msg = "All elements of window_strides must be positive, got {}."
        raise TypeError(msg.format(window_strides))
    if len(window_strides) != len(lhs_shape) - 2:
        msg = "{} window_strides has wrong length: expected {}, got {}."
        expected_length = len(lhs_shape) - 2
        raise TypeError(msg.format(name, expected_length, len(window_strides)))

def conv_shape_tuple(lhs_shape, rhs_shape, strides, pads, batch_group_count=1):
    """Compute the shape tuple of a conv given input shapes in canonical order."""
    if isinstance(pads, str):
        pads = lax.padtype_to_pads(lhs_shape[2:], rhs_shape[2:], strides, pads)
    if len(pads) != len(lhs_shape) - 2:
        msg = "Wrong number of explicit pads for convolution: expected {}, got {}."
        raise TypeError(msg.format(len(lhs_shape) - 2, len(pads)))

    lhs_padded = np.add(lhs_shape[2:], np.sum(np.array(pads).reshape(-1, 2),
                                              axis=1))

    if np.any(lhs_padded < 0):
        raise ValueError("Negative padding is larger than the size of the corresponding
dimension: "
                        f"got padding={pads} for lhs_shape[2:]={lhs_shape[2:]}")
    out_space = core.stride_shape(lhs_padded, rhs_shape[2:], strides)
    out_space = np.maximum(0, out_space)
    if batch_group_count > 1:
        assert lhs_shape[0] % batch_group_count == 0
        out_shape_0 = lhs_shape[0] // batch_group_count
    else:

```

```

out_shape_0 = lhs_shape[0]
out_shape = (out_shape_0, rhs_shape[0])
return tuple(out_shape + tuple(out_space))

def conv_general_shape_tuple(lhs_shape, rhs_shape, window_strides, padding,
                             dimension_numbers):
    lhs_perm, rhs_perm, out_perm = conv_general_permutations(dimension_numbers)
    lhs_trans = np.take(lhs_shape, lhs_perm)
    rhs_trans = np.take(rhs_shape, rhs_perm)
    out_trans = conv_shape_tuple(lhs_trans, rhs_trans, window_strides, padding)
    return tuple(np.take(out_trans, np.argsort(out_perm)))

def conv_transpose_shape_tuple(lhs_shape, rhs_shape, window_strides, padding,
                               dimension_numbers):
    lhs_perm, rhs_perm, out_perm = conv_general_permutations(dimension_numbers)
    lhs_trans = np.take(lhs_shape, lhs_perm)
    rhs_trans = np.take(rhs_shape, rhs_perm)
    if isinstance(padding, str):
        padding = [_conv_transpose_padding(k, s, padding)
                   for k, s in zip(rhs_trans[2:], window_strides)]
    padding = list(map(np.sum, padding))
    unpad_out_space = [(i-1) * s - k + 2
                       for i, k, s in zip(lhs_trans[2:],
                                           rhs_trans[2:],
                                           window_strides)]
    out_space = np.sum([unpad_out_space, padding], axis=0).tolist()
    out_trans = tuple((lhs_trans[0], rhs_trans[0]) + tuple(out_space))
    return tuple(np.take(out_trans, np.argsort(out_perm)))

def conv_dimension_numbers(lhs_shape, rhs_shape, dimension_numbers
                           ) -> ConvDimensionNumbers: \[docs\]
    """Converts convolution `dimension_numbers` to a `ConvDimensionNumbers`.

    Args:
        lhs_shape: tuple of nonnegative integers, shape of the convolution input.
        rhs_shape: tuple of nonnegative integers, shape of the convolution kernel.
        dimension_numbers: None or a tuple/list of strings or a ConvDimensionNumbers
            object following the convolution dimension number specification format in
            xla_client.py.

    Returns:
        A `ConvDimensionNumbers` object that represents `dimension_numbers` in the
        canonical form used by lax functions.
    """
    if isinstance(dimension_numbers, ConvDimensionNumbers):
        return dimension_numbers
    if len(lhs_shape) != len(rhs_shape):
        msg = "convolution requires lhs and rhs ndim to be equal, got {} and {}."
        raise TypeError(msg.format(len(lhs_shape), len(rhs_shape)))

    if dimension_numbers is None:
        iota = tuple(range(len(lhs_shape)))
        return ConvDimensionNumbers(iota, iota, iota)
    elif isinstance(dimension_numbers, (list, tuple)):
        if len(dimension_numbers) != 3:
            msg = "convolution dimension_numbers list/tuple must be length 3, got {}."
            raise TypeError(msg.format(len(dimension_numbers)))
        if not all(isinstance(elt, str) for elt in dimension_numbers):
            msg = "convolution dimension_numbers elements must be strings, got {}."
            raise TypeError(msg.format(tuple(map(type, dimension_numbers))))
        msg = ("convolution dimension_numbers[{}] must have len equal to the ndim "
              "of lhs and rhs, got {} for lhs shapes {} and {}.")
        for i, elt in enumerate(dimension_numbers):
            if len(elt) != len(lhs_shape):
                raise TypeError(msg.format(i, len(elt), lhs_shape, rhs_shape))

        lhs_spec, rhs_spec, out_spec = conv_general_permutations(dimension_numbers)
        return ConvDimensionNumbers(lhs_spec, rhs_spec, out_spec)
    else:
        msg = "convolution dimension_numbers must be tuple/list or None, got {}."
        raise TypeError(msg.format(type(dimension_numbers)))

def conv_general_permutations(dimension_numbers):
    """Utility for convolution dimension permutations relative to Conv HLO."""
    lhs_spec, rhs_spec, out_spec = dimension_numbers
    lhs_char, rhs_char, out_char = charpairs = ("N", "C"), ("O", "I"), ("N", "C")
    for i, (a, b) in enumerate(charpairs):
        if not dimension_numbers[i].count(a) == dimension_numbers[i].count(b) == 1:
            msg = ("convolution dimension_numbers[{}] must contain the characters "
                  "'{}' and '{}' exactly once, got {}.")
            raise TypeError(msg.format(i, a, b, dimension_numbers[i]))
        if len(dimension_numbers[i]) != len(set(dimension_numbers[i])):
            msg = ("convolution dimension_numbers[{}] cannot have duplicate "
                  "characters, got {}.")
            raise TypeError(msg.format(i, dimension_numbers[i]))
    if not (set(lhs_spec) - set(lhs_char) == set(rhs_spec) - set(rhs_char) ==

```

```

        set(out_spec) - set(out_char)):
    msg = ("convolution dimension_numbers elements must each have the same "
           "set of spatial characters, got {}.")
    raise TypeError(msg.format(dimension_numbers))

def getperm(spec, charpair):
    spatial = (i for i, c in enumerate(spec) if c not in charpair)
    if spec is not rhs_spec:
        spatial = sorted(spatial, key=lambda i: rhs_spec.index(spec[i]))
    return (spec.index(charpair[0]), spec.index(charpair[1])) + tuple(spatial)

lhs_perm, rhs_perm, out_perm = map(getperm, dimension_numbers, charpairs)
return lhs_perm, rhs_perm, out_perm

def _conv_general_proto(dimension_numbers):
    assert type(dimension_numbers) is ConvDimensionNumbers
    lhs_spec, rhs_spec, out_spec = dimension_numbers
    proto = xla_client.ConvolutionDimensionNumbers()
    proto.input_batch_dimension = lhs_spec[0]
    proto.input_feature_dimension = lhs_spec[1]
    proto.output_batch_dimension = out_spec[0]
    proto.output_feature_dimension = out_spec[1]
    proto.kernel_output_feature_dimension = rhs_spec[0]
    proto.kernel_input_feature_dimension = rhs_spec[1]
    proto.input_spatial_dimensions.extend(lhs_spec[2:])
    proto.kernel_spatial_dimensions.extend(rhs_spec[2:])
    proto.output_spatial_dimensions.extend(out_spec[2:])
    return proto

def _conv_general_vjp_lhs_padding(
    in_shape, window_dimensions, window_strides, out_shape, padding,
    lhs_dilation, rhs_dilation) -> List[Tuple[int, int]]:
    lhs_dilated_shape = lax._dilate_shape(in_shape, lhs_dilation)
    rhs_dilated_shape = lax._dilate_shape(window_dimensions, rhs_dilation)
    out_dilated_shape = lax._dilate_shape(out_shape, window_strides)
    pad_before = np.subtract(rhs_dilated_shape, [lo for lo, _ in padding]) - 1
    pad_after = (np.add(lhs_dilated_shape, rhs_dilated_shape) - 1
                 - out_dilated_shape - pad_before)
    return util.safe_zip(pad_before, pad_after)

def _conv_general_vjp_rhs_padding(
    in_shape, window_dimensions, window_strides, out_shape, padding,
    lhs_dilation, rhs_dilation):

    if len(in_shape) == 0: # 0D conv
        return []
    lhs_dilated_shape = lax._dilate_shape(in_shape, lhs_dilation)
    rhs_dilated_shape = lax._dilate_shape(window_dimensions, rhs_dilation)
    out_dilated_shape = lax._dilate_shape(out_shape, window_strides)
    pads_lo, _ = util.unzip2(padding)
    pads_from_lhs = core.diff_shape(out_dilated_shape, lhs_dilated_shape)
    pads_from_rhs = core.diff_shape(core.diff_shape(rhs_dilated_shape, pads_lo),
                                     (1,) * len(pads_lo))
    pads_hi = core.sum_shapes(pads_from_lhs, pads_from_rhs)
    return list(zip(pads_lo, pads_hi))

```

The JAX authors

Copyright 2020, The JAX Authors. NumPy and SciPy documentation are copyright the respective authors..