jax.numpy module

Contents

- jax.numpy.fft
- jax.numpy.linalg
- JAX DeviceArray

Implements the NumPy API, using the primitives in jax.lax.

While JAX tries to follow the NumPy API as closely as possible, sometimes JAX cannot follow NumPy exactly.

- Notably, since JAX arrays are immutable, NumPy APIs that mutate arrays in-place cannot be implemented in JAX. However, often JAX is able to provide an alternative API that is purely functional. For example, instead of in-place array updates (x[i] = y), JAX provides an alternative pure indexed update function x.at[i].set(y) (see ndarray.at).
- Relatedly, some NumPy functions often return views of arrays when possible (examples are <u>transpose()</u> and <u>reshape()</u>). JAX versions of such functions will return copies instead, although such are often optimized away by XLA when sequences of operations are compiled using <u>jax.jit()</u>.
- NumPy is very aggressive at promoting values to float64 type. JAX sometimes is less
 aggressive about type promotion (See <u>Type promotion semantics</u>).
- Some NumPy routines have data-dependent output shapes (examples include <u>unique()</u>)
 and <u>nonzero()</u>). Because the XLA compiler requires array shapes to be known at compile
 time, such operations are not compatible with JIT. For this reason, JAX adds an optional
 size argument to such functions which may be specified statically in order to use them
 with JIT.

Nearly all applicable NumPy functions are implemented in the jax.numpy namespace; they are listed below.

ndarray.at	Helper property for index update functionality.
<u>abs(x, /)</u>	Calculate the absolute value element- wise.
<pre>absolute(x, /)</pre>	Calculate the absolute value elementwise.
add(x1, x2, /)	Add arguments element-wise.
all(a[, axis, out, keepdims, where])	Test whether all array elements along a given axis evaluate to True.
allclose(a, b[, rtol, atol, equal_nan])	Returns True if two arrays are element-wise equal within a tolerance.
alltrue(a[, axis, out, keepdims, where])	Test whether all array elements along a given axis evaluate to True.
<pre>amax(a[, axis, out, keepdims, initial, where])</pre>	Return the maximum of an array or maximum along an axis.
amin(a[, axis, out, keepdims, initial, where])	Return the minimum of an array or minimum along an axis.
angle(z[, deg])	Return the angle of the complex argument.
any(a[, axis, out, keepdims, where])	Test whether any array element along a given axis evaluates to True.
append(arr, values[, axis])	Append values to the end of an array.
apply along axis(func1d, axis, arr, *args,)	Apply a function to 1-D slices along the given axis.
apply over axes(func, a, axes)	Apply a function repeatedly over multiple axes.

<pre>arange(start[, stop, step, dtype])</pre>	Return evenly spaced values within a given interval.
arccos(X, /)	Trigonometric inverse cosine, element-wise.
<pre>arccosh(x, /)</pre>	Inverse hyperbolic cosine, element- wise.
<pre>arcsin(X, /)</pre>	Inverse sine, element-wise.
<pre>arcsinh(x, /)</pre>	Inverse hyperbolic sine element-wise.
<pre>arctan(X, /)</pre>	Trigonometric inverse tangent, element-wise.
<u>arctan2(</u> x1, x2, /)	Element-wise arc tangent of x1/x2 choosing the quadrant correctly.
<pre>arctanh(x, /)</pre>	Inverse hyperbolic tangent element- wise.
argmax(a[, axis, out, keepdims])	Returns the indices of the maximum values along an axis.
<pre>argmin(a[, axis, out, keepdims])</pre>	Returns the indices of the minimum values along an axis.
<pre>argsort(a[, axis, kind, order])</pre>	Returns the indices that would sort an array.
argwhere(a, *[, size, fill_value])	Find the indices of array elements that are non-zero, grouped by element.
around(a[, decimals, out])	Evenly round to the given number of decimals.
<pre>array(object[, dtype, copy, order, ndmin])</pre>	Create an array.

<pre>array_equal(a1, a2[, equal_nan])</pre>	True if two arrays have the same shape and elements, False otherwise.
array equiv(a1, a2)	Returns True if input arrays are shape consistent and all elements equal.
<pre>array_repr(arr[, max_line_width, precision,])</pre>	Return the string representation of an array.
<pre>array_split(ary, indices_or_sections[, axis])</pre>	Split an array into multiple sub-arrays.
<pre>array_str(a[, max_line_width, precision,])</pre>	Return a string representation of the data in an array.
asarray(a[, dtype, order])	Convert the input to an array.
atleast 1d(*arys)	Convert inputs to arrays with at least one dimension.
atleast 2d(*arys)	View inputs as arrays with at least two dimensions.
atleast 3d(*arys)	View inputs as arrays with at least three dimensions.
average()	Compute the weighted average along the specified axis.
<pre>bartlett(M)</pre>	Return the Bartlett window.
<pre>bincount(x[, weights, minlength, length])</pre>	Count number of occurrences of each value in array of non-negative ints.
<pre>bitwise and(x1, x2, /)</pre>	Compute the bit-wise AND of two arrays element-wise.
<pre>bitwise not(X, /)</pre>	Compute bit-wise inversion, or bit-wise NOT, element-wise.

Compute the bit-wise OR of two arrays element-wise.
Compute the bit-wise XOR of two arrays element-wise.
Return the Blackman window.
Assemble an nd-array from nested lists of blocks.
Broadcast any number of arrays against each other.
Broadcast the input shapes into a single shape.
Broadcast an array to a new shape.
Concatenate slices, scalars and array- like objects along the last axis.
Returns True if cast between data types can occur according to the casting rule.
Return the cube-root of an array, element-wise.
alias of jax.numpy.complex128
Return the ceiling of the input, element-wise.
Abstract base class of all character

<pre>choose(a, choices[, out, mode])</pre>	Construct an array from an index array and a list of arrays to choose from.
clip(a[, a_min, a_max, out])	Clip (limit) the values in an array.
column stack(tup)	Stack 1-D arrays as columns into a 2-D array.
complex	alias of jax.numpy.complex128
<pre>complex128(X)</pre>	
<pre>complex64(X)</pre>	
<pre>complexfloating()</pre>	Abstract base class of all complex number scalar types that are made up of floating-point numbers.
ComplexWarning	The warning raised when casting a complex dtype to a real dtype.
<pre>compress(condition, a[, axis, out])</pre>	Return selected slices of an array along given axis.
<pre>concatenate(arrays[, axis, dtype])</pre>	Join a sequence of arrays along an existing axis.
conj(x, /)	Return the complex conjugate, element-wise.
<pre>conjugate(x, /)</pre>	Return the complex conjugate, element-wise.
convolve(a, v[, mode, precision])	Returns the discrete, linear convolution of two one-dimensional sequences.
<u>сору</u> (а[, order])	Return an array copy of the given object.

<pre>copysign(x1, x2, /)</pre>	Change the sign of x1 to that of x2, element-wise.
<pre>corrcoef(x[, y, rowvar])</pre>	Return Pearson product-moment correlation coefficients.
<pre>correlate(a, v[, mode, precision])</pre>	Cross-correlation of two 1-dimensional sequences.
<u>cos</u> (x, /)	Cosine element-wise.
<u>cosh</u> (x, /)	Hyperbolic cosine, element-wise.
<pre>count nonzero(a[, axis, keepdims])</pre>	Counts the number of non-zero values in the array a.
cov(m[, y, rowvar, bias, ddof, fweights,])	Estimate a covariance matrix, given data and weights.
cross(a, b[, axisa, axisb, axisc, axis])	Return the cross product of two (arrays of) vectors.
csingle	alias of jax.numpy.complex64
<pre>cumprod(a[, axis, dtype, out])</pre>	Return the cumulative product of elements along a given axis.
<pre>cumproduct(a[, axis, dtype, out])</pre>	Return the cumulative product of elements along a given axis.
<pre>cumsum(a[, axis, dtype, out])</pre>	Return the cumulative sum of the elements along a given axis.
deg2rad(x, /)	Convert angles from degrees to radians.
degrees(X, /)	Convert angles from radians to degrees.

<pre>delete(arr, obj[, axis])</pre>	Return a new array with sub-arrays along an axis deleted.
<pre>diag(v[, k])</pre>	Extract a diagonal or construct a diagonal array.
<pre>diag indices(n[, ndim])</pre>	Return the indices to access the main diagonal of an array.
diag indices from(arr)	Return the indices to access the main diagonal of an n-dimensional array.
<pre>diagflat(V[, k])</pre>	Create a two-dimensional array with the flattened input as a diagonal.
diagonal(a[, offset, axis1, axis2])	Return specified diagonals.
<pre>diff(a[, n, axis, prepend, append])</pre>	Calculate the n-th discrete difference along the given axis.
<pre>digitize(x, bins[, right])</pre>	Return the indices of the bins to which each value in input array belongs.
<u>divide(</u> x1, x2, /)	Divide arguments element-wise.
<u>divmod</u> (x1, x2, /)	Return element-wise quotient and remainder simultaneously.
<pre>dot(a, b, *[, precision])</pre>	Dot product of two arrays.
double	alias of jax.numpy.float64
<pre>dsplit(ary, indices_or_sections)</pre>	Split array into multiple sub-arrays along the 3rd axis (depth).
dstack(tup[, dtype])	Stack arrays in sequence depth wise (along third axis).
dtype(dtype[, align, copy])	Create a data type object.

<pre>ediff1d(ary[, to_end, to_begin])</pre>	The differences between consecutive elements of an array.
einsum(*operands[, out, optimize,])	Evaluates the Einstein summation convention on the operands.
<pre>einsum path(subscripts, *operands[, optimize])</pre>	Evaluates the lowest cost contraction order for an einsum expression by
<pre>empty(shape[, dtype])</pre>	Return a new array of given shape and type, without initializing entries.
<pre>empty like(prototype[, dtype, shape])</pre>	Return a new array with the same shape and type as a given array.
<u>equal(</u> x1, x2, /)	Return (x1 == x2) element-wise.
<u>exp(</u> x, /)	Calculate the exponential of all elements in the input array.
<u>exp2(</u> X, /)	Calculate $2^{**}p$ for all p in the input array.
expand dims(a, axis)	Expand the shape of an array.
<u>expm1</u> (X, /)	Calculate $\exp(x) - 1$ for all elements in the array.
<pre>expm1(X, /) extract(condition, arr)</pre>	
	in the array. Return the elements of an array that
extract(condition, arr)	in the array. Return the elements of an array that satisfy some condition. Return a 2-D array with ones on the

<pre>fix(x[, out])</pre>	Round to nearest integer towards zero.
flatnonzero(a, *[, size, fill_value])	Return indices that are non-zero in the flattened version of a.
<pre>flexible()</pre>	Abstract base class of all scalar types without predefined length.
<pre>flip(m[, axis])</pre>	Reverse the order of elements in an array along the given axis.
<pre>fliplr(m)</pre>	Reverse the order of elements along axis 1 (left/right).
<u>flipud</u> (m)	Reverse the order of elements along axis 0 (up/down).
<u>float</u>	alias of jax.numpy.float64
<pre>float power(x1, x2, /)</pre>	First array elements raised to powers from second array, element-wise.
float16(X)	
float32(X)	
float64(X)	
<pre>floating()</pre>	Abstract base class of all floating-point scalar types.
<pre>floor(x, /)</pre>	Return the floor of the input, element- wise.
<pre>floor divide(x1, x2, /)</pre>	Return the largest integer smaller or equal to the division of the inputs.
<u>fmax(</u> X1, X2)	Element-wise maximum of array elements.

<u>fmin(</u> x1, x2)	Element-wise minimum of array elements.
<u>fmod(</u> x1, x2, /)	Returns the element-wise remainder of division.
<pre>frexp(X, /)</pre>	Decompose the elements of x into mantissa and twos exponent.
<pre>frombuffer(buffer[, dtype, count, offset])</pre>	Interpret a buffer as a 1-dimensional array.
<pre>fromfile(*args, **kwargs)</pre>	Unimplemented JAX wrapper for jnp.fromfile.
<pre>fromfunction(function, shape, *[, dtype])</pre>	Construct an array by executing a function over each coordinate.
<pre>fromiter(*args, **kwargs)</pre>	Unimplemented JAX wrapper for jnp.fromiter.
<pre>fromstring(string[, dtype, count])</pre>	A new 1-D array initialized from text data in a string.
from dlpack(X)	Create a NumPy array from an object implementing thedlpack
<pre>full(shape, fill_value[, dtype])</pre>	Return a new array of given shape and type, filled with <i>fill_value</i> .
<pre>full like(a, fill_value[, dtype, shape])</pre>	Return a full array with the same shape and type as a given array.
gcd(x1, x2)	Returns the greatest common divisor of x1 and x2
generic()	Base class for numpy scalar types.
<pre>geomspace(start, stop[, num, endpoint,])</pre>	Return numbers spaced evenly on a log scale (a geometric progression).

<pre>get_printoptions()</pre>	Return the current print options.
<pre>gradient(f, *varargs[, axis, edge_order])</pre>	Return the gradient of an N-dimensional array.
<pre>greater(x1, x2, /)</pre>	Return the truth value of (x1 > x2) element-wise.
<pre>greater equal(x1, x2, /)</pre>	Return the truth value of $(x1 \ge x2)$ element-wise.
<pre>hamming(M)</pre>	Return the Hamming window.
hanning(M)	Return the Hanning window.
<pre>heaviside(x1, x2, /)</pre>	Compute the Heaviside step function.
<pre>histogram(a[, bins, range, weights, density])</pre>	Compute the histogram of a dataset.
<pre>histogram bin edges(a[, bins, range, weights])</pre>	Function to calculate only the edges of the bins used by the <i>histogram</i>
<pre>histogram2d(x, y[, bins, range, weights,])</pre>	Compute the bi-dimensional histogram of two data samples.
<pre>histogramdd(sample[, bins, range, weights,])</pre>	Compute the multidimensional histogram of some data.
<pre>hsplit(ary, indices_or_sections)</pre>	Split an array into multiple sub-arrays horizontally (column-wise).
<pre>hstack(tup[, dtype])</pre>	Stack arrays in sequence horizontally (column wise).
<u>hypot(</u> x1, x2, /)	Given the "legs" of a right triangle, return its hypotenuse.
<u>i0(</u> X)	Modified Bessel function of the first kind, order 0.

<pre>identity(n[, dtype])</pre>	Return the identity array.
<u>iinfo</u> (type)	Machine limits for integer types.
<pre>imag(val, /)</pre>	Return the imaginary part of the complex argument.
<u>in1d</u> (ar1, ar2[, assume_unique, invert])	Test whether each element of a 1-D array is also present in a second array.
<u>index exp</u>	A nicer way to build up index tuples for arrays.
indices()	Return an array representing the indices of a grid.
<pre>inexact()</pre>	Abstract base class of all numeric scalar types with a (potentially) inexact representation of the values in its range, such as floating-point numbers.
<u>inner</u> (a, b, *[, precision])	Inner product of two arrays.
<pre>insert(arr, obj, values[, axis])</pre>	Insert values along the given axis before the given indices.
<u>int</u>	alias of jax.numpy.int64
<u>int16(</u> X)	
<u>int32(</u> X)	
<u>int64(</u> X)	
int8(X)	
<pre>integer()</pre>	Abstract base class of all integer scalar types.

<pre>interp(x, xp, fp[, left, right, period])</pre>	One-dimensional linear interpolation for monotonically increasing sample points.
<pre>intersect1d(ar1, ar2[, assume_unique,])</pre>	Find the intersection of two arrays.
<pre>invert(x, /)</pre>	Compute bit-wise inversion, or bit-wise NOT, element-wise.
<u>isclose(a, b[, rtol, atol, equal_nan])</u>	Returns a boolean array where two arrays are element-wise equal within a
<u>iscomplex</u> (X)	Returns a bool array, where True if input element is complex.
<u>iscomplexobj</u> (X)	Check for a complex type or an array of complex numbers.
<pre>isfinite(x, /)</pre>	Test element-wise for finiteness (not infinity and not Not a Number).
isin(element, test_elements[,])	Calculates element in test_elements, broadcasting over <i>element</i> only.
<pre>isinf(x, /)</pre>	Test element-wise for positive or negative infinity.
<u>isnan(</u> X, /)	Test element-wise for NaN and return result as a boolean array.
<pre>isneginf(x, /[, out])</pre>	Test element-wise for negative infinity, return result as bool array.
<pre>isposinf(x, /[, out])</pre>	Test element-wise for positive infinity, return result as bool array.
<u>isreal</u> (X)	Returns a bool array, where True if input element is real.

<u>isrealobj(</u> X)	Return True if x is a not complex type or an array of complex numbers.
<u>isscalar</u> (element)	Returns True if the type of <i>element</i> is a scalar type.
<u>issubdtype</u> (arg1, arg2)	Returns True if first argument is a typecode lower/equal in type hierarchy.
<u>issubsctype</u> (arg1, arg2)	Determine if the first argument is a subclass of the second argument.
<pre>iterable(y)</pre>	Check whether or not an object can be iterated over.
ix (*args)	Construct an open mesh from multiple sequences.
kaiser(M, beta)	Return the Kaiser window.
kron(a, b)	Kronecker product of two arrays.
<u>lcm(</u> x1, x2)	Returns the lowest common multiple of x1 and x2
<u>ldexp(</u> x1, x2, /)	Returns x1 * 2**x2, element-wise.
<pre>left shift(x1, x2, /)</pre>	Shift the bits of an integer to the left.
<u>less(</u> x1, x2, /)	Return the truth value of (x1 < x2) element-wise.
<pre>less equal(x1, x2, /)</pre>	Return the truth value of (x1 <= x2) element-wise.
<u>lexsort</u> (keys[, axis])	Perform an indirect stable sort using a sequence of keys.

<u>linspace</u> ()	Return evenly spaced numbers over a specified interval.
load(*args, **kwargs)	Load arrays or pickled objects from .npy, .npz or pickled files.
<u>log(</u> x, /)	Natural logarithm, element-wise.
<u>log10(</u> X, /)	Return the base 10 logarithm of the input array, element-wise.
<u>log1p(</u> X, /)	Return the natural logarithm of one plus the input array, element-wise.
<u>log2(</u> x, /)	Base-2 logarithm of x.
logaddexp(x1, x2, /)	Logarithm of the sum of exponentiations of the inputs.
logaddexp2(X1, X2, /)	Logarithm of the sum of exponentiations of the inputs in base-2.
logical and (*args)	Compute the truth value of x1 AND x2 element-wise.
logical not(*args)	Compute the truth value of NOT x element-wise.
logical or(*args)	Compute the truth value of x1 OR x2 element-wise.
<u>logical xor</u> (*args)	Compute the truth value of x1 XOR x2, element-wise.
logspace(start, stop[, num, endpoint, base,])	Return numbers spaced evenly on a log scale.
mask indices(*args, **kwargs)	Return the indices to access (n, n) arrays, given a masking function.

<pre>matmul(a, b, *[, precision])</pre>	Matrix product of two arrays.
<pre>max(a[, axis, out, keepdims, initial, where])</pre>	Return the maximum of an array or maximum along an axis.
<u>maximum</u> (x1, x2, /)	Element-wise maximum of array elements.
mean(a[, axis, dtype, out, keepdims, where])	Compute the arithmetic mean along the specified axis.
<pre>median(a[, axis, out, overwrite_input, keepdims])</pre>	Compute the median along the specified axis.
meshgrid(*xi[, copy, sparse, indexing])	Return coordinate matrices from coordinate vectors.
mgrid	Return dense multi-dimensional "meshgrid".
min(a[, axis, out, keepdims, initial, where])	Return the minimum of an array or minimum along an axis.
<pre>minimum(x1, x2, /)</pre>	Element-wise minimum of array elements.
mod(x1, x2, /)	Returns the element-wise remainder of division.
<u>modf</u> (x, /[, out])	Return the fractional and integral parts of an array, element-wise.
moveaxis(a, source, destination)	Move axes of an array to new positions.
msort(a)	Return a copy of an array sorted along the first axis.
<u>multiply(</u> x1, x2, /)	Multiply arguments element-wise.

<pre>nan to num(x[, copy, nan, posinf, neginf])</pre>	Replace NaN with zero and infinity with large finite numbers (default
nanargmax(a[, axis, out, keepdims])	Return the indices of the maximum values in the specified axis ignoring
<pre>nanargmin(a[, axis, out, keepdims])</pre>	Return the indices of the minimum values in the specified axis ignoring
<pre>nancumprod(a[, axis, dtype, out])</pre>	Return the cumulative product of array elements over a given axis treating Not a
nancumsum(a[, axis, dtype, out])	Return the cumulative sum of array elements over a given axis treating Not a
nanmax(a[, axis, out, keepdims, initial, where])	Return the maximum of an array or maximum along an axis, ignoring any
nanmean(a[, axis, dtype, out, keepdims, where])	Compute the arithmetic mean along the specified axis, ignoring NaNs.
<pre>nanmedian(a[, axis, out, overwrite_input,])</pre>	Compute the median along the specified axis, while ignoring NaNs.
nanmin(a[, axis, out, keepdims, initial, where])	Return minimum of an array or minimum along an axis, ignoring any NaNs.
nanpercentile(a, q[, axis, out,])	Compute the qth percentile of the data along the specified axis,
nanprod(a[, axis, dtype, out, keepdims,])	Return the product of array elements over a given axis treating Not a
nanquantile(a, q[, axis, out,])	Compute the qth quantile of the data along the specified axis,

<pre>nanstd(a[, axis, dtype, out, ddof,])</pre>	Compute the standard deviation along the specified axis, while
nansum(a[, axis, dtype, out, keepdims,])	Return the sum of array elements over a given axis treating Not a
nanvar(a[, axis, dtype, out, ddof,])	Compute the variance along the specified axis, while ignoring NaNs.
<u>ndarray</u>	alias of jax.Array
ndim(a)	Return the number of dimensions of an array.
<pre>negative(x, /)</pre>	Numerical negative, element-wise.
<pre>nextafter(x1, x2, /)</pre>	Return the next floating-point value after x1 towards x2, element-wise.
nonzero(a, *[, size, fill_value])	Return the indices of the elements that are non-zero.
<pre>not equal(x1, x2, /)</pre>	Return (x1 != x2) element-wise.
number()	Abstract base class of all numeric scalar types.
<u>object</u>	Any Python object.
<u>ogrid</u>	Return open multi-dimensional "meshgrid".
ones(shape[, dtype])	Return a new array of given shape and type, filled with ones.
ones like(a[, dtype, shape])	Return an array of ones with the same shape and type as a given array.
outer(a, b[, out])	Compute the outer product of two vectors.

<pre>packbits(a[, axis, bitorder])</pre>	Packs the elements of a binary-valued array into bits in a uint8 array.
<pre>pad(array, pad_width[, mode])</pre>	Pad an array.
<pre>partition(a, kth[, axis])</pre>	Return a partitioned copy of an array.
<pre>percentile(a, q[, axis, out,])</pre>	Compute the q-th percentile of the data along the specified axis.
<pre>piecewise(x, condlist, funclist, *args, **kw)</pre>	Evaluate a piecewise-defined function.
place(*args, **kwargs)	Change elements of an array based on conditional and input values.
<pre>poly(seq_of_zeros)</pre>	Find the coefficients of a polynomial with the given sequence of roots.
polyadd(a1, a2)	Find the sum of two polynomials.
<pre>polyder(p[, m])</pre>	Return the derivative of the specified order of a polynomial.
<pre>polydiv(u, v, *[, trim_leading_zeros])</pre>	Returns the quotient and remainder of polynomial division.
<pre>polyfit(x, y, deg[, rcond, full, w, cov])</pre>	Least squares polynomial fit.
<pre>polyint(p[, m, k])</pre>	Return an antiderivative (indefinite integral) of a polynomial.
<pre>polymul(a1, a2, *[, trim_leading_zeros])</pre>	Find the product of two polynomials.
polysub(a1, a2)	Difference (subtraction) of two polynomials.
<pre>polyval(p, x, *[, unroll])</pre>	Evaluate a polynomial at specific values.
<pre>positive(X, /)</pre>	Numerical positive, element-wise.

<u>power</u> (x1, x2, /)	First array elements raised to powers from second array, element-wise.
<u>printoptions</u> (*args, **kwargs)	Context manager for setting print options.
<pre>prod(a[, axis, dtype, out, keepdims,])</pre>	Return the product of array elements over a given axis.
<pre>product(a[, axis, dtype, out, keepdims,])</pre>	Return the product of array elements over a given axis.
<pre>promote types(a, b)</pre>	Returns the type to which a binary operation should cast its arguments.
<pre>ptp(a[, axis, out, keepdims])</pre>	Range of values (maximum - minimum) along an axis.
<u>put</u> (*args, **kwargs)	Replaces specified elements of an array with given values.
<pre>quantile(a, q[, axis, out, overwrite_input,])</pre>	Compute the q-th quantile of the data along the specified axis.
<u>r_</u>	Concatenate slices, scalars and array- like objects along the first axis.
rad2deg(x, /)	Convert angles from radians to degrees.
radians(X, /)	Convert angles from degrees to radians.
ravel(a[, order])	Return a contiguous flattened array.
<pre>ravel multi index(multi_index, dims[, mode,])</pre>	Converts a tuple of index arrays into an array of flat
real(val, /)	Return the real part of the complex argument.

<pre>reciprocal(X, /)</pre>	Return the reciprocal of the argument, element-wise.
<pre>remainder(x1, x2, /)</pre>	Returns the element-wise remainder of division.
<pre>repeat(a, repeats[, axis, total_repeat_length])</pre>	Repeat elements of an array.
reshape(a, newshape[, order])	Gives a new shape to an array without changing its data.
resize(a, new_shape)	Return a new array with the specified shape.
result type(*args)	Returns the type that results from applying the NumPy
<pre>right shift(x1, x2, /)</pre>	Shift the bits of an integer to the right.
<u>rint(</u> X, /)	Round elements of the array to the nearest integer.
roll(a, shift[, axis])	Roll array elements along a given axis.
rollaxis(a, axis[, start])	Roll the specified axis backwards, until it lies in a given position.
<pre>roots(p, *[, strip_zeros])</pre>	Return the roots of a polynomial with coefficients given in p.
<u>rot90(</u> m[, k, axes])	Rotate an array by 90 degrees in the plane specified by axes.
round(a[, decimals, out])	Evenly round to the given number of decimals.
round (a[, decimals, out])	Evenly round to the given number of decimals.

<pre>row stack(tup[, dtype])</pre>	Stack arrays in sequence vertically (row wise).
<u>S</u>	A nicer way to build up index tuples for arrays.
<pre>save(file, arr[, allow_pickle, fix_imports])</pre>	Save an array to a binary file in NumPy .npy format.
<pre>savez(file, *args, **kwds)</pre>	Save several arrays into a single file in uncompressed .npz format.
searchsorted(a, v[, side, sorter, method])	Find indices where elements should be inserted to maintain order.
<pre>select(condlist, choicelist[, default])</pre>	Return an array drawn from elements in choicelist, depending on conditions.
<pre>set printoptions([precision, threshold,])</pre>	Set printing options.
<pre>setdiff1d(ar1, ar2[, assume_unique, size,])</pre>	Find the set difference of two arrays.
<pre>setxor1d(ar1, ar2[, assume_unique])</pre>	Find the set exclusive-or of two arrays.
shape(a)	Return the shape of an array.
<pre>sign(X, /)</pre>	Returns an element-wise indication of the sign of a number.
<pre>signbit(X, /)</pre>	Returns element-wise True where signbit is set (less than zero).
<pre>signedinteger()</pre>	Abstract base class of all signed integer scalar types.
<u>sin(</u> X, /)	Trigonometric sine, element-wise.
<u>sinc(</u> X, /)	Return the normalized sinc function.
<u>single</u>	alias of jax.numpy.float32

sinh(x, /)	Hyperbolic sine, element-wise.
<pre>size(a[, axis])</pre>	Return the number of elements along
	a given axis.
<pre>sometrue(a[, axis, out, keepdims, where])</pre>	Test whether any array element along
	a given axis evaluates to True.
<pre>sort(a[, axis, kind, order])</pre>	Return a sorted copy of an array.
<pre>sort complex(a)</pre>	Sort a complex array using the real
	part first, then the imaginary part.
<pre>split(ary, indices_or_sections[, axis])</pre>	Split an array into multiple sub-arrays
	as views into <i>ary</i> .
<u>sqrt(</u> X, /)	Return the non-negative square-root
	of an array, element-wise.
<u>square(</u> X, /)	Return the element-wise square of the
	input.
squeeze(a[, axis])	Remove axes of length one from a.
stack(arrays[, axis, out, dtype])	Join a sequence of arrays along a new
	axis.
<pre>std(a[, axis, dtype, out, ddof, keepdims, where])</pre>	Compute the standard deviation along
	the specified axis.
<pre>subtract(x1, x2, /)</pre>	Subtract arguments, element-wise.
<pre>sum(a[, axis, dtype, out, keepdims,])</pre>	Sum of array elements over a given
	axis.
swapaxes(a, axis1, axis2)	Interchange two axes of an array.
take(a, indices[, axis, out, mode,])	Take elements from an array along an
	axis.

take along axis(arr, indices, axis[, mode])	Take values from the input array by matching 1d index and data slices.
<u>tan(</u> x, /)	Compute tangent element-wise.
tanh(x, /)	Compute hyperbolic tangent element- wise.
tensordot(a, b[, axes, precision])	Compute tensor dot product along specified axes.
tile(A, reps)	Construct an array by repeating A the number of times given by reps.
trace(a[, offset, axis1, axis2, dtype, out])	Return the sum along diagonals of the array.
transpose(a[, axes])	Returns an array with axes transposed.
trapz(y[, x, dx, axis])	Integrate along the given axis using
<u>trapz</u> (y[, x, ux, axis])	the composite trapezoidal rule.
tri(N[, M, k, dtype])	
	the composite trapezoidal rule. An array with ones at and below the
tri(N[, M, k, dtype])	the composite trapezoidal rule. An array with ones at and below the given diagonal and zeros elsewhere.
<pre>tri(N[, M, k, dtype]) tril(m[, k])</pre>	the composite trapezoidal rule. An array with ones at and below the given diagonal and zeros elsewhere. Lower triangle of an array. Return the indices for the lower-
<pre>tri(N[, M, k, dtype]) tril(m[, k]) tril indices(*args, **kwargs)</pre>	the composite trapezoidal rule. An array with ones at and below the given diagonal and zeros elsewhere. Lower triangle of an array. Return the indices for the lower-triangle of an (n, m) array. Return the indices for the lower-
<pre>tri(N[, M, k, dtype]) tril(m[, k]) tril indices(*args, **kwargs) tril indices from(arr[, k])</pre>	the composite trapezoidal rule. An array with ones at and below the given diagonal and zeros elsewhere. Lower triangle of an array. Return the indices for the lower-triangle of an (n, m) array. Return the indices for the lower-triangle of arr. Trim the leading and/or trailing zeros

triu indices(*args, **kwargs)	Return the indices for the upper- triangle of an (n, m) array.
triu indices from(arr[, k])	Return the indices for the upper- triangle of arr.
<pre>true divide(X1, X2, /)</pre>	Divide arguments element-wise.
trunc(X)	Return the truncated value of the input, element-wise.
<u>uint</u>	alias of jax.numpy.uint64
<u>uint16(</u> X)	
<u>uint32(</u> X)	
uint64(X)	
<pre>uint8(X)</pre>	
union1d(ar1, ar2, *[, size, fill_value])	Find the union of two arrays.
unique(ar[, return_index, return_inverse,])	Find the unique elements of an array.
unpackbits(a[, axis, count, bitorder])	Unpacks elements of a uint8 array into a binary-valued output array.
unravel index(indices, shape)	Converts a flat index or array of flat indices into a tuple
<pre>unsignedinteger()</pre>	Abstract base class of all unsigned integer scalar types.
unwrap(p[, discont, axis, period])	Unwrap by taking the complement of large deltas with respect to the period.
vander(x[, N, increasing])	Generate a Vandermonde matrix.

var(a[, axis, dtype, out, ddof, keepdims, where])	Compute the variance along the specified axis.
<pre>vdot(a, b, *[, precision])</pre>	Return the dot product of two vectors.
<pre>vectorize(pyfunc, *[, excluded, signature])</pre>	Define a vectorized function with broadcasting.
<pre>vsplit(ary, indices_or_sections)</pre>	Split an array into multiple sub-arrays vertically (row-wise).
vstack(tup[, dtype])	Stack arrays in sequence vertically (row wise).
where()	Return elements chosen from x or y depending on <i>condition</i> .
zeros(shape[, dtype])	Return a new array of given shape and type, filled with zeros.
zeros like(a[, dtype, shape])	Return an array of zeros with the same shape and type as a given array.

jax.numpy.fft

<pre>fft(a[, n, axis, norm])</pre>	Compute the one-dimensional discrete Fourier Transform.
<pre>fft2(a[, s, axes, norm])</pre>	Compute the 2-dimensional discrete Fourier Transform.
fftfreq(n[, d, dtype])	Return the Discrete Fourier Transform sample frequencies.
<pre>fftn(a[, s, axes, norm])</pre>	Compute the N-dimensional discrete Fourier Transform.
<pre>fftshift(x[, axes])</pre>	Shift the zero-frequency component to the center of the spectrum.
<pre>hfft(a[, n, axis, norm])</pre>	Compute the FFT of a signal that has Hermitian symmetry, i.e., a real
ifft(a[, n, axis, norm])	Compute the one-dimensional inverse discrete Fourier Transform.
ifft2(a[, s, axes, norm])	Compute the 2-dimensional inverse discrete Fourier Transform.
<u>ifftn(a[, s, axes, norm])</u>	Compute the N-dimensional inverse discrete Fourier Transform.
<pre>ifftshift(x[, axes])</pre>	The inverse of fftshift.
ihfft(a[, n, axis, norm])	Compute the inverse FFT of a signal that has Hermitian symmetry.
<pre>irfft(a[, n, axis, norm])</pre>	Computes the inverse of rfft.
<u>irfft2(a[, s, axes, norm])</u>	Computes the inverse of rfft2.
<u>irfftn(a[, s, axes, norm])</u>	Computes the inverse of rfftn.
rfft(a[, n, axis, norm])	Compute the one-dimensional discrete Fourier Transform for real input.
rfft2(a[, s, axes, norm])	Compute the 2-dimensional FFT of a real array.

<pre>rfftfreq(n[, d, dtype])</pre>	Return the Discrete Fourier Transform sample frequencies
<pre>rfftn(a[, s, axes, norm])</pre>	Compute the N-dimensional discrete Fourier Transform for real input.

jax.numpy.linalg

<u>cholesky</u> (a)	Cholesky decomposition.
cond(x[, p])	Compute the condition number of a matrix.
<u>det</u> (a)	Compute the determinant of an array.
<u>eig</u> (a)	Compute the eigenvalues and right eigenvectors of a square array.
<pre>eigh(a[, UPLO, symmetrize_input])</pre>	Return the eigenvalues and eigenvectors of a complex Hermitian
<u>eigvals</u> (a)	Compute the eigenvalues of a general matrix.
eigvalsh(a[, UPLO])	Compute the eigenvalues of a complex Hermitian or real symmetric matrix.
<u>inv</u> (a)	Compute the (multiplicative) inverse of a matrix.
<pre>lstsq(a, b[, rcond, numpy_resid])</pre>	Return the least-squares solution to a linear matrix equation.
<pre>matrix power(a, n)</pre>	Raise a square matrix to the (integer) power n .
<pre>matrix_rank(M[, tol])</pre>	Return matrix rank of array using SVD method
<pre>multi dot(arrays, *[, precision])</pre>	Compute the dot product of two or more arrays in a single function call,
<pre>norm(x[, ord, axis, keepdims])</pre>	Matrix or vector norm.
pinv(a[, rcond, hermitian])	Compute the (Moore-Penrose) pseudo-inverse of a matrix.
qr()	Compute the qr factorization of a matrix.
slogdet(a, *[, method])	Compute the sign and (natural) logarithm of the determinant of an array.

solve(a, b)	Solve a linear matrix equation, or system of linear scalar equations.
svd()	Singular Value Decomposition.
tensorinv(a[, ind])	Compute the 'inverse' of an N-dimensional array.
tensorsolve(a, b[, axes])	Solve the tensor equation $a \times = b$ for x .

JAX DeviceArray

The JAX <u>DeviceArray</u> is the core array object in JAX: you can think of it as the equivalent of a <u>numpy.ndarray</u> backed by a memory buffer on a single device. Like <u>numpy.ndarray</u>, most users will not need to instantiate <u>DeviceArray</u> objects manually, but rather will create them via <u>jax.numpy</u> functions like <u>array()</u>, <u>arange()</u>, <u>linspace()</u>, and others listed above.

Copying and Serialization

DeviceArray` objects are designed to work seamlessly with Python standard library tools where appropriate.

With the built-in <u>copy</u> module, when <u>copy.copy()</u> or <u>copy.deepcopy()</u> encounder a <u>DeviceArray</u>, it is equivalent to calling the <u>copy()</u> method, which will create a copy of the buffer on the same device as the original array. This will work correctly within traced/JIT-compiled code, though copy operations may be elided by the compiler in this context.

When the built-in <u>pickle</u> module encounters a <u>DeviceArray</u>, it will be serialized via a compact bit representation in a similar manner to pickled <u>numpy.ndarray</u> objects. When unpickled, the result will be a new <u>DeviceArray</u> object on the default device. This is because in general, pickling and unpickling may take place in different runtime environments, and there is no general way to map the device IDs of one runtime to the device IDs of another. If <u>pickle</u> is used in traced/JIT-compiled code, it will result in a <u>ConcretizationTypeError</u>.

Class Reference

jax.numpy.**DeviceArray**

alias of jaxlib.xla extension.DeviceArrayBase

class jaxlib.xla_extension.DeviceArrayBase

class jaxlib.xla_extension.DeviceArray

property T: jax.Array

Returns an array with axes transposed.

LAX-backend implementation of numpy.transpose().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

For a 1-D array, this returns an unchanged view of the original array, as a transposed vector is simply the same vector. To convert a 1-D array into a 2-D column vector, an additional dimension must be added, e.g., np.atleast2d(a).T achieves this, as does a[:, np.newaxis]. For a 2-D array, this is the standard matrix transpose. For an n-D array, if axes are given, their order indicates how the axes are permuted (see Examples). If axes are not provided, then transpose(a).shape == a.shape[::-1].

Parameters:

- **a** (*array_like*) Input array.
- axes (<u>tuple</u> or list of ints, optional) If specified, it must be a tuple or list which contains a permutation of [0,1,...,N-1] where N is the number of axes of a. The i'th axis of the returned array will correspond to the axis numbered axes[i] of the input. If not specified, defaults to range(a.ndim)[::-1], which reverses the order of the axes.

Returns: p - a with its axes permuted. A view is returned whenever possible.

Return type: ndarray

all(axis=None, out=None, keepdims=False, *, where=None)

Test whether all array elements along a given axis evaluate to True.

LAX-backend implementation of numpy.all().

Original docstring below.

Parameters:

- a (array_like) Input array or object that can be converted to an array.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along which a logical AND reduction is performed. The default (axis=None) is to perform a logical AND over all the dimensions of the input array. axis may be negative, in which case it counts from the last to the first axis.

• keepdims (bool, optional) -

If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *all* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- where (array_like of bool, optional) Elements to include in checking for all *True* values. See ~numpy.ufunc.reduce for details.
- out (<u>None</u>) -

Returns: all – A new boolean or array is returned unless *out* is specified, in

which case a reference to out is returned.

Return type: ndarray, bool

any(axis=None, out=None, keepdims=False, *, where=None)

Test whether any array element along a given axis evaluates to True.

LAX-backend implementation of numpy.any().

Original docstring below.

Returns single boolean if axis is None

Parameters:

- a (array_like) Input array or object that can be converted to an array.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along which a logical OR reduction is performed. The default (axis=None) is to perform a logical OR over all the dimensions of the input array. axis may be negative, in which case it counts from the last to the first axis.
- **keepdims** (<u>bool</u>, optional) –

If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *any* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- where (array_like of bool, optional) Elements to include in checking for any True values. See ~numpy.ufunc.reduce for details.
- out (None) -

Returns: any – A new boolean or *ndarray* is returned unless *out* is specified,

in which case a reference to *out* is returned.

Return type: <u>bool</u> or ndarray

argmax(axis=None, out=None, keepdims=None)

Returns the indices of the maximum values along an axis.

LAX-backend implementation of numpy.argmax().

Original docstring below.

Parameters: • a (array_like) – Input array.

 axis (<u>int</u>, optional) – By default, the index is into the flattened array, otherwise along the specified axis.

keepdims (bool, optional) – If this is set to True, the axes which
are reduced are left in the result as dimensions with size one.
With this option, the result will broadcast correctly against the
array.

Returns: in

index_array – Array of indices into the array. It has the same shape as *a.shape* with the dimension along *axis* removed. If *keepdims* is set to True, then the size of *axis* will be 1 with the resulting array having same shape as *a.shape*.

Return type: ndarray of ints

argmin(axis=None, out=None, keepdims=None)

Returns the indices of the minimum values along an axis.

LAX-backend implementation of numpy.argmin().

Original docstring below.

Parameters:

- **a** (array_like) Input array.
- axis (<u>int</u>, optional) By default, the index is into the flattened array, otherwise along the specified axis.
- keepdims (bool, optional) If this is set to True, the axes which
 are reduced are left in the result as dimensions with size one.
 With this option, the result will broadcast correctly against the
 array.

Returns: index_array – Array of indices into the array. It has the same shape

as *a.shape* with the dimension along *axis* removed. If *keepdims* is set to True, then the size of *axis* will be 1 with the resulting array

having same shape as a.shape.

Return type: ndarray of ints

argpartition(**kwargs)

Perform an indirect partition along the given axis using the

LAX-backend implementation of numpy.argpartition().

* This function is not yet implemented by jax.numpy, and will raise NotImplementedError *

Original docstring below.

algorithm specified by the *kind* keyword. It returns an array of indices of the same shape as *a* that index data along the given axis in partitioned order.

New in version 1.8.0.

Parameters:

- a (array_like) Array to sort.
- **kth** (<u>int</u> or sequence of ints) –

Element index to partition by. The k-th element will be in its final sorted position and all smaller elements will be moved before it and all larger elements behind it. The order of all elements in the partitions is undefined. If provided with a sequence of k-th it will partition all of them into their sorted position at once.

Deprecated since version 1.22.0: Passing booleans as index is deprecated.

- axis (<u>int</u> or None, optional) Axis along which to sort. The default is -1 (the last axis). If None, the flattened array is used.
- kind ({'introselect'}, optional) Selection algorithm. Default is 'introselect'
- order (<u>str</u> or list of str, optional) When a is an array with fields
 defined, this argument specifies which fields to compare first,
 second, etc. A single field can be specified as a string, and not all

fields need be specified, but unspecified fields will still be used, in the order in which they come up in the dtype, to break ties.

Returns: index_array – Array of indices that partition a along the specified

axis. If a is one-dimensional, a[index_array] yields a partitioned a.

More generally, np.take_along_axis(a, index_array, axis=axis) always yields the partitioned a, irrespective of

dimensionality.

Return type: ndarray, int

argsort(axis=- 1, kind='stable', order=None)

Returns the indices that would sort an array.

LAX-backend implementation of numpy.argsort().

Only kind='stable' is supported. Other kind values will produce a warning and be treated as if they were 'stable'.

Original docstring below.

Perform an indirect sort along the given axis using the algorithm specified by the *kind* keyword. It returns an array of indices of the same shape as *a* that index data along the given axis in sorted order.

Parameters:

- **a** (array_like) Array to sort.
- axis (<u>int</u> or None, optional) Axis along which to sort. The default is -1 (the last axis). If None, the flattened array is used.
- **kind** ({'quicksort', 'mergesort', 'heapsort', 'stable'}, optional) Sorting algorithm. The default is 'quicksort'. Note that both 'stable' and 'mergesort' use timsort under the covers and, in general, the actual implementation will vary with data type. The 'mergesort' option is retained for backwards compatibility.
 - Changed in version 1.15.0.: The 'stable' option was added.
- order (<u>str</u> or list of str, optional) When a is an array with fields defined, this argument specifies which fields to compare first, second, etc. A single field can be specified as a string, and not all fields need be specified, but unspecified fields will still be used, in the order in which they come up in the dtype, to break ties.

Returns: index_array – Array of indices that sort *a* along the specified *axis*.

If a is one-dimensional, a[index_array] yields a sorted a. More generally, np.take_along_axis(a, index_array, axis=axis)

always yields the sorted a, irrespective of dimensionality.

Return type: ndarray, <u>int</u>

astype(dtype)

Copy the array and cast to a specified dtype.

This is implemeted via <u>jax.lax.convert element type()</u>, which may have slightly different behavior than <u>numpy.ndarray.astype()</u> in some cases. In particular, the details of float-to-int and int-to-float casts are implementation dependent.

Parameters: • arr (<u>Union[Array, ndarray, bool, number, bool, int, float</u>,

complex]) -

dtype (<u>Union[Any</u>, <u>str</u>, <u>dtype</u>, <u>SupportsDType]</u>) –

Return type: Array

property at

Helper property for index update functionality.

The at property provides a functionally pure equivalent of in-place array modifications.

In particular:

Alternate syntax	Equivalent In-place expression
x = x.at[idx].set(y)	x[idx] = y
x = x.at[idx].add(y)	x[idx] += y
<pre>x = x.at[idx].multiply(y)</pre>	x[idx] *= y
<pre>x = x.at[idx].divide(y)</pre>	x[idx] /= y
x = x.at[idx].power(y)	x[idx] **= y
<pre>x = x.at[idx].min(y)</pre>	x[idx] = minimum(x[idx], y)
x = x.at[idx].max(y)	x[idx] = maximum(x[idx], y)
<pre>x = x.at[idx].apply(ufunc)</pre>	ufunc.at(x, idx)
x = x.at[idx].get()	x = x[idx]

None of the x.at expressions modify the original x; instead they return a modified copy of x. However, inside a $\underline{\mathtt{jit}()}$ compiled function, expressions like x = x.at[idx].set(y) are guaranteed to be applied in-place.

Unlike NumPy in-place operations such as x[idx] += y, if multiple indices refer to the same location, all updates will be applied (NumPy would only apply the last update, rather than applying all updates.) The order in which conflicting updates are applied is implementation-defined and may be nondeterministic (e.g., due to concurrency on some hardware platforms).

By default, JAX assumes that all indices are in-bounds. Alternative out-of-bound index semantics can be specified via the mode parameter (see below).

- mode (<u>str</u>)
 - Specify out-of-bound indexing mode. Options are:
 - "promise_in_bounds": (default) The user promises that indices are in bounds. No additional checking will be performed. In practice, this means that out-of-bounds indices in get() will be clipped, and out-of-bounds indices in set(), add(), etc. will be dropped.
 - "clip": clamp out of bounds indices into valid range.

- "drop": ignore out-of-bound indices.
- "fill": alias for "drop". For get(), the optional fill_value argument specifies the value that will be returned.
 - See <u>jax.lax.GatherScatterMode</u> for more details.
- indices_are_sorted (<u>bool</u>) If True, the implementation will
 assume that the indices passed to at[] are sorted in ascending
 order, which can lead to more efficient execution on some
 backends.
- unique_indices (bool) If True, the implementation will assume
 that the indices passed to at[] are unique, which can result in
 more efficient execution on some backends.
- **fill_value** (*Any*) Only applies to the get() method: the fill value to return for out-of-bounds slices when *mode* is 'fill'. Ignored otherwise. Defaults to NaN for inexact types, the largest negative value for signed types, the largest positive value for unsigned types, and True for booleans.

Examples

```
>>> x = inp.arange(5.0)
Array([0., 1., 2., 3., 4.], dtype=float32)
>>> x.at[2].add(10)
Array([ 0., 1., 12., 3., 4.], dtype=float32)
>>> x.at[10].add(10) # out-of-bounds indices are ignored
Array([0., 1., 2., 3., 4.], dtype=float32)
>>> x.at[20].add(10, mode='clip')
Array([ 0., 1., 2., 3., 14.], dtype=float32)
>>> x.at[2].get()
Array(2., dtype=float32)
>>> x.at[20].get() # out-of-bounds indices clipped
Array(4., dtype=float32)
>>> x.at[20].get(mode='fill') # out-of-bounds indices filled with NaN
Array(nan, dtype=float32)
>>> x.at[20].get(mode='fill', fill_value=-1) # custom fill value
Array(-1., dtype=float32)
```

block_until_ready()

(self: xla::PyBuffer::pyobject) -> StatusOr[xla::PyBuffer::pyobject]

broadcast(sizes)

Broadcasts an array, adding new leading dimensions

```
    operand (<u>Union[Array, ndarray, bool, number, bool, int, float, complex]</u>) – an array
```

sizes (<u>Sequence[int]</u>) – a sequence of integers, giving the sizes
of new leading dimensions to add to the front of the array.

Return type: Array

Returns: An array containing the result.

See also

jax.lax.broadcast_in_dim : add new dimensions at any location in the array shape.

broadcast_in_dim(shape, broadcast_dimensions)

Wraps XLA's **BroadcastInDim** operator.

Parameters: • operand (<u>Union[Array</u>, <u>ndarray</u>, <u>bool</u>, <u>number</u>, <u>bool</u>, <u>int</u>,

float, complex]) - an array

• **shape** (Sequence[Union[int, Any]]) – the shape of the target

array

• **broadcast_dimensions** (<u>Sequence[int]</u>) – to which dimension in

the target shape each dimension of the operand shape

corresponds to

Return type: Array

Returns: An array containing the result.

See also

jax.lax.broadcast: simpler interface to add new leading dimensions.

choose(choices, out=None, mode='raise')

Construct an array from an index array and a list of arrays to choose from.

LAX-backend implementation of numpy.choose().

Original docstring below.

First of all, if confused or uncertain, definitely look at the Examples - in its full generality, this function is less simple than it might seem from the following code description (below ndi = numpy.lib.index_tricks):

```
np.choose(a,c) == np.array([c[a[I]][I] for I in ndi.ndindex(a.shape)]).
```

But this omits some subtleties. Here is a fully general summary:

Given an "index" array (a) of integers and a sequence of n arrays (*choices*), a and each choice array are first broadcast, as necessary, to arrays of a common shape; calling these Ba and Bchoices[i], i = 0,...,n-1 we have that, necessarily, Ba.shape == Bchoices[i].shape for each i. Then, a new array with shape Ba.shape is created as follows:

- if mode='raise' (the default), then, first of all, each element of a (and thus Ba) must be in the range [0, n-1]; now, suppose that i (in that range) is the value at the (j0, j1, ..., jm) position in Ba then the value at the same position in the new array is the value in Bchoices[i] at that same position;
- if mode='wrap', values in a (and thus Ba) may be any (signed) integer; modular arithmetic is used to map integers outside the range [0, n-1] back into that range; and then the new array is constructed as above;
- if mode='clip', values in a (and thus Ba) may be any (signed) integer; negative integers are mapped to 0; values greater than n-1 are mapped to n-1; and then the new array is constructed as above.

Parameters:

- a (int array) This array must contain integers in [0, n-1],
 where n is the number of choices, unless mode=wrap or
 mode=clip, in which cases any integers are permissible.
- choices (sequence of arrays) Choice arrays. a and all of the choices must be broadcastable to the same shape. If choices is itself an array (not recommended), then its outermost dimension (i.e., the one corresponding to choices.shape[0]) is taken as defining the "sequence".
- mode ({'raise' (default), 'wrap', 'clip'}, optional) –
 Specifies how indices outside [0, n-1] will be treated:
 - 'raise': an exception is raised
 - 'wrap': value becomes value mod n
 - 'clip': values < 0 are mapped to 0, values > n-1 are mapped to n-1
- out (<u>None</u>) -

Returns: merged_array – The merged result.

Return type: array

clone()

(self: xla::PyBuffer::pyobject) -> xla::PyBuffer::pyobject

conj()

Return the complex conjugate, element-wise.

LAX-backend implementation of numpy.conjugate().

Original docstring below.

The complex conjugate of a complex number is obtained by changing the sign of its imaginary part.

Parameters: x (*array_like*) – Input value.

Returns: y – The complex conjugate of x, with same dtype as y. This is a

scalar if x is a scalar.

Return type: ndarray

conjugate()

Return the complex conjugate, element-wise.

LAX-backend implementation of numpy.conjugate().

Original docstring below.

The complex conjugate of a complex number is obtained by changing the sign of its imaginary part.

Parameters: x (*array_like*) – Input value.

Returns: y – The complex conjugate of x, with same dtype as y. This is a

scalar if x is a scalar.

Return type: ndarray

copy(order=None)

Return an array copy of the given object.

LAX-backend implementation of numpy.copy().

This function will create arrays on JAX's default device. For control of the device placement of data, see <u>jax.device put()</u>. More information is available in the JAX FAQ at <u>Controlling data and computation placement on devices</u> (full FAQ at https://jax.readthedocs.io/en/latest/faq.html).

Original docstring below.

Parameters: • a (array_like) – Input data.

order ({'C', 'F', 'A', 'K'}, optional) – Controls the memory layout of the copy. 'C' means C-order, 'F' means F-order, 'A' means 'F' if a is Fortran contiguous, 'C' otherwise. 'K' means match the layout of a as closely as possible. (Note that this function and ndarray.copy() are very similar, but have different default values for their order= arguments.)

Returns: arr – Array interpretation of *a*.

Return type: ndarray

copy_to_device()

(self: xla::PyBuffer::pyobject, arg0: jaxlib.xla extension.Device) -> StatusOr[object]

copy_to_host_async()

(self: xla::PyBuffer::pyobject) -> Status

copy_to_remote_device()

(self: xla::PyBuffer::pyobject, arg0: bytes) -> Tuple[Status, bool]

cumprod(axis=None, dtype=None, out=None)

Return the cumulative product of elements along a given axis.

LAX-backend implementation of numpy.cumprod().

Original docstring below.

Parameters:

- a (array_like) Input array.
- axis (<u>int</u>, optional) Axis along which the cumulative product is computed. By default the input is flattened.
- dtype (<u>dtype</u>, optional) Type of the returned array, as well as of
 the accumulator in which the elements are multiplied. If dtype is
 not specified, it defaults to the dtype of a, unless a has an integer
 dtype with a precision less than that of the default platform
 integer. In that case, the default platform integer is used instead.
- out (<u>None</u>) -

Returns: cumprod – A new array holding the result is returned unless *out* is

specified, in which case a reference to out is returned.

Return type: ndarray

cumsum(axis=None, dtype=None, out=None)

Return the cumulative sum of the elements along a given axis.

LAX-backend implementation of numpy.cumsum().

Original docstring below.

Parameters:

- a (array_like) Input array.
- axis (<u>int</u>, optional) Axis along which the cumulative sum is computed. The default (None) is to compute the cumsum over the flattened array.
- dtype (<u>dtype</u>, optional) Type of the returned array and of the accumulator in which the elements are summed. If *dtype* is not specified, it defaults to the dtype of a, unless a has an integer dtype with a precision less than that of the default platform integer. In that case, the default platform integer is used.
- out (<u>None</u>) -

Returns:

cumsum_along_axis – A new array holding the result is returned unless *out* is specified, in which case a reference to *out* is returned. The result has the same size as *a*, and the same shape as *a* if *axis* is not None or *a* is a 1-d array.

Return type: ndarray.

delete()

(self: xla::PyBuffer::pyobject) -> None

device()

(self: xla::PyBuffer::pyobject) -> jaxlib.xla extension.Device

diagonal(offset=0, axis1=0, axis2=1)

Return specified diagonals.

LAX-backend implementation of numpy.diagonal().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

If a is 2-D, returns the diagonal of a with the given offset, i.e., the collection of elements of the form a[i, i+offset]. If a has more than two dimensions, then the axes specified by axis1 and axis2 are used to determine the 2-D sub-array whose diagonal is returned. The shape of the resulting array can be determined by removing axis1 and axis2 and appending an index to the right equal to the size of the resulting diagonals.

In versions of NumPy prior to 1.7, this function always returned a new, independent array containing a copy of the values in the diagonal.

In NumPy 1.7 and 1.8, it continues to return a copy of the diagonal, but depending on this fact is deprecated. Writing to the resulting array continues to work as it used to, but a FutureWarning is issued.

Starting in NumPy 1.9 it returns a read-only view on the original array. Attempting to write to the resulting array will produce an error.

In some future release, it will return a read/write view and writing to the returned array will alter your original array. The returned array will have the same type as the input array.

If you don't write to the array returned by this function, then you can just ignore all of the above.

If you depend on the current behavior, then we suggest copying the returned array explicitly, i.e., use np.diagonal(a).copy() instead of just np.diagonal(a). This will work with both past and future versions of NumPy.

Parameters:

- **a** (*array_like*) Array from which the diagonals are taken.
- offset (<u>int</u>, optional) Offset of the diagonal from the main diagonal. Can be positive or negative. Defaults to main diagonal (0).
- axis1 (<u>int</u>, optional) Axis to be used as the first axis of the 2-D sub-arrays from which the diagonals should be taken. Defaults to first axis (0).
- axis2 (<u>int</u>, optional) Axis to be used as the second axis of the 2-D sub-arrays from which the diagonals should be taken. Defaults to second axis (1).

Returns:

array_of_diagonals – If *a* is 2-D, then a 1-D array containing the diagonal and of the same type as *a* is returned unless *a* is a *matrix*, in which case a 1-D array rather than a (2-D) *matrix* is returned in order to maintain backward compatibility.

If a.ndim > 2, then the dimensions specified by axis1 and axis2 are removed, and a new axis inserted at the end corresponding to the diagonal.

Return type: ndarray

dot(b, *, precision=None)

Dot product of two arrays. Specifically,

LAX-backend implementation of numpy.dot().

In addition to the original NumPy arguments listed below, also supports precision for extra control over matrix-multiplication precision on supported devices. precision may be set to None, which means default precision for the backend, a <u>Precision</u> enum value (Precision.DEFAULT, Precision.HIGH or Precision.HIGHEST) or a tuple of two <u>Precision</u> enums indicating separate precision for each argument.

Original docstring below.

- If both *a* and *b* are 1-D arrays, it is inner product of vectors (without complex conjugation).
- If both a and b are 2-D arrays, it is matrix multiplication, but using matmul() or a
 @ b is preferred.
- If either a or b is 0-D (scalar), it is equivalent to multiply() and using numpy.multiply(a, b) or a * b is preferred.
- If a is an N-D array and b is a 1-D array, it is a sum product over the last axis of a
 and b.
- If a is an N-D array and b is an M-D array (where M>=2), it is a sum product over the last axis of a and the second-to-last axis of b:

```
dot(a, b)[i,j,k,m] = sum(a[i,j,:] * b[k,:,m])
```

It uses an optimized BLAS library when possible (see *numpy.linalg*).

• a (array like) – First argument.

• **b** (array_like) – Second argument.

Returns: output – Returns the dot product of a and b. If a and b are both

scalars or both 1-D arrays then a scalar is returned; otherwise an

array is returned. If *out* is given, then it is returned.

Return type: ndarray

flatten(order='C')

Return a contiguous flattened array.

LAX-backend implementation of numpy.ravel().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

A 1-D array, containing the elements of the input, is returned. A copy is made only if needed.

As of NumPy 1.10, the returned array will have the same type as the input array. (for example, a masked array will be returned for a masked array input)

Parameters:

- a (array_like) Input array. The elements in a are read in the order specified by order, and packed as a 1-D array.
- **order** ({'C', 'F', 'A', 'K'}, optional) The elements of a are read using this index order. 'C' means to index the elements in rowmajor, C-style order, with the last axis index changing fastest, back to the first axis index changing slowest. 'F' means to index the elements in column-major, Fortran-style order, with the first index changing fastest, and the last index changing slowest. Note that the 'C' and 'F' options take no account of the memory layout of the underlying array, and only refer to the order of axis indexing. 'A' means to read the elements in Fortran-like index order if a is Fortran contiguous in memory, C-like order otherwise. 'K' means to read the elements in the order they occur in memory, except for reversing the data when strides are negative. By default, 'C' index order is used.

Returns:

 \mathbf{y} – y is an array of the same subtype as a, with shape (a.size,).

Note that matrices are special cased for backward compatibility, if a

is a matrix, then y is a 1-D ndarray.

Return type: array like

property imag: jax.Array

Return the imaginary part of the complex argument.

LAX-backend implementation of numpy.imag().

Original docstring below.

Parameters: val (array_like) – Input array.

Returns: out – The imaginary component of the complex argument. If *val* is

real, the type of *val* is used for the output. If *val* has complex

elements, the returned type is float.

Return type: ndarray or scalar

is_deleted()

(self: xla::PyBuffer::pyobject) -> bool

is_known_ready()

(self: xla::PyBuffer::pyobject) -> StatusOr[bool]

is_ready()

(self: xla::PyBuffer::pyobject) -> StatusOr[bool]

max(axis=None, out=None, keepdims=False, initial=None, where=None)

Return the maximum of an array or maximum along an axis.

LAX-backend implementation of numpy.amax().

Original docstring below.

Parameters:

- a (array_like) Input data.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along which to operate. By default, flattened input is used.
- keepdims (bool, optional) –
 If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.
 If the default value is passed, then keepdims will not be passed

through to the *amax* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- initial (scalar, optional) The minimum value of an output element. Must be present to allow computation on empty slice.
 See ~numpy.ufunc.reduce for details.
- where (array_like of bool, optional) Elements to compare for the maximum. See ~numpy.ufunc.reduce for details.
- out (<u>None</u>) -

Returns:

amax - Maximum of a. If axis is None, the result is a scalar value.
If axis is an int, the result is an array of dimension a.ndim - 1. If
axis is a tuple, the result is an array of dimension a.ndim len(axis).

Return type: ndarray or scalar

mean(axis=None, dtype=None, out=None, keepdims=False, *,
where=None)

Compute the arithmetic mean along the specified axis.

LAX-backend implementation of numpy.mean().

Original docstring below.

Returns the average of the array elements. The average is taken over the flattened array by default, otherwise over the specified axis. *float64* intermediate and return values are used for integer inputs.

Parameters:

- a (array_like) Array containing numbers whose mean is desired. If a is not an array, a conversion is attempted.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along
 which the means are computed. The default is to compute the
 mean of the flattened array.
- dtype (data-type, optional) Type to use in computing the mean.
 For integer inputs, the default is float64; for floating point inputs, it is the same as the input dtype.
- keepdims (<u>bool</u>, optional) –
 If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.
 If the default value is passed, then keepdims will not be passed through to the magn method of sub classes of pdarray because.
 - through to the *mean* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.
- where (array_like of bool, optional) Elements to include in the mean. See ~numpy.ufunc.reduce for details.
- out (None) -

Returns:

m – If *out=None*, returns a new array containing the mean values, otherwise a reference to the output array is returned.

Return type: ndarray, see dtype parameter above

min(axis=None, out=None, keepdims=False, initial=None, where=None)

Return the minimum of an array or minimum along an axis.

LAX-backend implementation of numpy.amin().

Original docstring below.

- a (array_like) Input data.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along which to operate. By default, flattened input is used.
- keepdims (bool, optional) –
 If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *amin* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- initial (scalar, optional) The maximum value of an output element. Must be present to allow computation on empty slice. See ~numpy.ufunc.reduce for details.
- where (array_like of bool, optional) Elements to compare for the minimum. See ~numpv.ufunc.reduce for details.
- out (<u>None</u>) -

Returns: amin – Minimum of a. If axis is None, the result is a scalar value. If

axis is an int, the result is an array of dimension a.ndim - 1. If axis is a tuple, the result is an array of dimension a.ndim - len(axis).

Return type: ndarray or scalar

nonzero(*, size=None, fill_value=None)

Return the indices of the elements that are non-zero.

LAX-backend implementation of numpy.nonzero().

Because the size of the output of nonzero is data-dependent, the function is not typically compatible with JIT. The JAX version adds the optional size argument which must be specified statically for jnp.nonzero to be used within some of JAX's transformations.

Original docstring below.

Returns a tuple of arrays, one for each dimension of *a*, containing the indices of the non-zero elements in that dimension. The values in *a* are always tested and returned in row-major, C-style order.

To group the indices by element, rather than dimension, use *argwhere*, which returns a row for each non-zero element.

Note

When called on a zero-d array or scalar, nonzero(a) is treated as $nonzero(atleast_1d(a))$.

① *Deprecated since version 1.17.0:* Use *atleast_1d* explicitly if this behavior is deliberate.

Parameters:

- a (array_like) Input array.
- **size** (*int*, *optional*) If specified, the indices of the first size True elements will be returned. If there are fewer unique elements than size indicates, the return value will be padded with fill value.
- fill_value (array_like, optional) When size is specified and there are fewer than the indicated number of elements, the remaining elements will be filled with fill_value, which defaults to zero.

Returns: tuple_of_arrays – Indices of elements that are non-zero.

Return type: <u>tuple</u>

on_device_size_in_bytes()

(self: xla::PyBuffer::pyobject) -> StatusOr[int]

platform()

(self: xla::PyBuffer::pyobject) -> str

prod(axis=None, dtype=None, out=None, keepdims=False, initial=None,
where=None, promote_integers=True)

Return the product of array elements over a given axis.

LAX-backend implementation of numpy.prod().

Original docstring below.

- a (array_like) Input data.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along
 which a product is performed. The default, axis=None, will
 calculate the product of all the elements in the input array. If axis
 is negative it counts from the last to the first axis.
- dtype (<u>dtype</u>, optional) The type of the returned array, as well as of the accumulator in which the elements are multiplied. The dtype of a is used by default unless a has an integer dtype of less precision than the default platform integer. In that case, if a is signed then the platform integer is used while if a is unsigned then an unsigned integer of the same precision as the platform integer is used.
- **keepdims** (<u>bool</u>, optional) –

If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *prod* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- initial (scalar, optional) The starting value for this product. See ~numpy.ufunc.reduce for details.
- where (array like of bool, optional) Elements to include in the product. See ~numpy.ufunc.reduce for details.
- **promote_integers** (<u>bool</u>, default=True) If True, then integer inputs will be promoted to the widest available integer dtype. following numpy's behavior. If False, the result will have the same dtype as the input. promote_integers is ignored if dtype is specified.
- out (<u>None</u>) -

product_along_axis - An array shaped as a but with the specified Returns:

axis removed. Returns a reference to out if specified.

Return type: ndarray, see *dtype* parameter above.

ptp(axis=None, out=None, keepdims=False)

Range of values (maximum - minimum) along an axis.

LAX-backend implementation of numpy.ptp().

Original docstring below.

The name of the function comes from the acronym for 'peak to peak'.



Warning

ptp preserves the data type of the array. This means the return value for an input of signed integers with n bits (e.g. np.int8, np.int16, etc) is also a signed integer with n bits. In that case, peak-to-peak values greater than 2**(n-1)-1 will be returned as negative values. An example with a work-around is shown below.

Parameters:a (array_like) – Input values.

- axis (None or <u>int</u> or tuple of ints, optional) Axis along which to
 find the peaks. By default, flatten the array. axis may be negative,
 in which case it counts from the last to the first axis.
- keepdims (bool, optional) -

If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *ptp* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

• out (<u>None</u>) -

Returns: ptp – The range of a given array - scalar if array is one-dimensional

or a new array holding the result along the given axis

Return type: ndarray or scalar

ravel(order='C')

Return a contiguous flattened array.

LAX-backend implementation of numpy.ravel().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

A 1-D array, containing the elements of the input, is returned. A copy is made only if needed.

As of NumPy 1.10, the returned array will have the same type as the input array. (for example, a masked array will be returned for a masked array input)

- a (array_like) Input array. The elements in a are read in the order specified by order, and packed as a 1-D array.
- order ({'C', 'F', 'A', 'K'}, optional) The elements of a are read using this index order. 'C' means to index the elements in row-major, C-style order, with the last axis index changing fastest, back to the first axis index changing slowest. 'F' means to index the elements in column-major, Fortran-style order, with the first index changing fastest, and the last index changing slowest. Note that the 'C' and 'F' options take no account of the memory layout of the underlying array, and only refer to the order of axis

indexing. 'A' means to read the elements in Fortran-like index order if *a* is Fortran *contiguous* in memory, C-like order otherwise.

'K' means to read the elements in the order they occur in

memory, except for reversing the data when strides are negative.

By default, 'C' index order is used.

Returns: y - y is an array of the same subtype as a, with shape (a.size,).

Note that matrices are special cased for backward compatibility, if a

is a matrix, then y is a 1-D ndarray.

Return type: array_like

property real: jax.Array

Return the real part of the complex argument.

LAX-backend implementation of numpy.real().

Original docstring below.

Parameters: val (array_like) – Input array.

Returns: out – The real component of the complex argument. If *val* is real,

the type of val is used for the output. If val has complex elements,

the returned type is float.

Return type: ndarray or scalar

repeat (repeats, axis=None, *, total_repeat_length=None)

Repeat elements of an array.

LAX-backend implementation of numpy.repeat().

JAX adds the optional *total_repeat_length* parameter which specifies the total number of repeat, and defaults to sum(repeats). It must be specified for repeat to be compilable. If *sum(repeats)* is larger than the specified *total_repeat_length* the remaining values will be discarded. In the case of *sum(repeats)* being smaller than the specified target length, the final value will be repeated.

Original docstring below.

• a (array_like) – Input array.

- **repeats** (*int* or array of ints) The number of repetitions for each element. *repeats* is broadcasted to fit the shape of the given axis.
- axis (<u>int</u>, optional) The axis along which to repeat values. By default, use the flattened input array, and return a flat output array.

total_repeat_length (<u>Optional[int]</u>) -

Returns: repeated array – Output array which has the same shape as a,

except along the given axis.

Return type: ndarray

round(decimals=0, out=None)

Evenly round to the given number of decimals.

LAX-backend implementation of numpy.around().

Original docstring below.

Parameters: • a (array_like) – Input data.

 decimals (<u>int</u>, optional) – Number of decimal places to round to (default: 0). If decimals is negative, it specifies the number of

positions to the left of the decimal point.

Returns: rounded_array – An array of the same type as *a*, containing the

rounded values. Unless out was specified, a new array is created. A

reference to the result is returned.

The real and imaginary parts of complex numbers are rounded

separately. The result of rounding a float is a float.

Return type: ndarray

References

1 "Lecture Notes on the Status of IEEE 754", William Kahan, https://people.eecs.berkeley.edu/~wkahan/ieee754status/IEEE754.PDF

Parameters: out (None) –

searchsorted(v, side='left', sorter=None, *, method='scan')

Find indices where elements should be inserted to maintain order.

LAX-backend implementation of numpy.searchsorted().

Original docstring below.

Find the indices into a sorted array a such that, if the corresponding elements in v were inserted before the indices, the order of a would be preserved.

Assuming that a is sorted:

side	returned index <i>i</i> satisfies
left	a[i-1] < v <= a[i]
right	a[i-1] <= v < a[i]

Parameters:

- a (1-D array_like) Input array. If sorter is None, then it must be sorted in ascending order, otherwise sorter must be an array of indices that sort it.
- v (array_like) Values to insert into a.
- **side** (*{'left', 'right'}, optional*) If 'left', the index of the first suitable location found is given. If 'right', return the last such index. If there is no suitable index, return either 0 or N (where N is the length of *a*).
- method (<u>str</u>) One of 'scan' (default) or 'sort'. Controls the
 method used by the implementation; 'scan' tends to be more
 performant on CPU (particularly when a is very large), while 'sort'
 is often more performant on accelerator backends like GPU and
 TPU (particularly when v is very large).
- sorter (None) -

Returns: indices – Array of insertion points with the same shape as v, or an

integer if v is a scalar.

Return type: <u>int</u> or array of ints

sort(axis=- 1, kind='quicksort', order=None)

Return a sorted copy of an array.

LAX-backend implementation of numpy.sort().

Original docstring below.

Parameters: • **a** (array_like) – Array to be sorted.

- axis (<u>int</u> or None, optional) Axis along which to sort. If None, the array is flattened before sorting. The default is -1, which sorts along the last axis.
- kind ({'quicksort', 'mergesort', 'heapsort', 'stable'}, optional) –
 Sorting algorithm. The default is 'quicksort'. Note that both 'stable' and 'mergesort' use timsort or radix sort under the covers and, in general, the actual implementation will vary with data

type. The 'mergesort' option is retained for backwards compatibility.

① Changed in version 1.15.0.: The 'stable' option was added.

order (str or list of str, optional) – When a is an array with fields defined, this argument specifies which fields to compare first, second, etc. A single field can be specified as a string, and not all fields need be specified, but unspecified fields will still be used, in the order in which they come up in the dtype, to break ties.

Returns: sorted_array – Array of the same type and shape as *a*.

Return type: ndarray

split(indices_or_sections, axis=0)

Split an array into multiple sub-arrays as views into ary.

LAX-backend implementation of numpy.split().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

Parameters:

- **ary** (*ndarray*) Array to be divided into sub-arrays.
- indices_or_sections (<u>int</u> or 1-D array) –
 If indices_or_sections is an integer, N, the array will be divided into N equal arrays along axis. If such a split is not possible, an error is raised.

If *indices_or_sections* is a 1-D array of sorted integers, the entries indicate where along *axis* the array is split. For example, [2, 3] would, for axis=0, result in

- ary[:2]
- ary[2:3]
- ary[3:]

If an index exceeds the dimension of the array along *axis*, an empty sub-array is returned correspondingly.

• axis (<u>int</u>, optional) – The axis along which to split, default is 0.

Returns: sub-arrays – A list of sub-arrays as views into *ary*.

Return type: list of ndarrays

squeeze(axis=None)

Remove axes of length one from a.

LAX-backend implementation of numpy.squeeze().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

Parameters: • a (array_like) – Input data.

• axis (None or int or tuple of ints, optional) –

Returns: squeezed – The input array, but with all or a subset of the

dimensions of length 1 removed. This is always a itself or a view into a. Note that if all axes are squeezed, the result is a 0d array

and not a scalar.

Return type: ndarray

std(axis=None, dtype=None, out=None, ddof=0, keepdims=False, *,
where=None)

Compute the standard deviation along the specified axis.

LAX-backend implementation of numpy.std().

Original docstring below.

Returns the standard deviation, a measure of the spread of a distribution, of the array elements. The standard deviation is computed for the flattened array by default, otherwise over the specified axis.

- **a** (*array_like*) Calculate the standard deviation of these values.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along
 which the standard deviation is computed. The default is to
 compute the standard deviation of the flattened array.
- dtype (<u>dtype</u>, optional) Type to use in computing the standard deviation. For arrays of integer type the default is float64, for arrays of float types it is the same as the array type.
- ddof (<u>int</u>, optional) Means Delta Degrees of Freedom. The
 divisor used in calculations is N ddof, where N represents the
 number of elements. By default ddof is zero.
- **keepdims** (<u>bool</u>, optional) –

If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *std* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- where (array_like of bool, optional) Elements to include in the standard deviation. See ~numpy.ufunc.reduce for details.
- out (None) -

Returns:

standard_deviation – If *out* is None, return a new array containing the standard deviation, otherwise return a reference to the output array.

Return type: ndarray, see dtype parameter above.

sum(axis=None, dtype=None, out=None, keepdims=False, initial=None,
where=None, promote_integers=True)

Sum of array elements over a given axis.

LAX-backend implementation of numpy.sum().

Original docstring below.

- **a** (array_like) Elements to sum.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along
 which a sum is performed. The default, axis=None, will sum all of
 the elements of the input array. If axis is negative it counts from
 the last to the first axis.
- dtype (<u>dtype</u>, optional) The type of the returned array and of
 the accumulator in which the elements are summed. The dtype of
 a is used by default unless a has an integer dtype of less
 precision than the default platform integer. In that case, if a is
 signed then the platform integer is used while if a is unsigned
 then an unsigned integer of the same precision as the platform
 integer is used.
- keepdims (bool, optional) –
 If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.

If the default value is passed, then *keepdims* will not be passed through to the *sum* method of sub-classes of *ndarray*, however any non-default value will be. If the sub-class' method does not implement *keepdims* any exceptions will be raised.

- initial (scalar, optional) Starting value for the sum. See
 ~numpy.ufunc.reduce for details.
- where (array_like of bool, optional) Elements to include in the sum. See ~numpy.ufunc.reduce for details.
- promote_integers (<u>bool</u>, default=True) If True, then integer inputs will be promoted to the widest available integer dtype, following numpy's behavior. If False, the result will have the same dtype as the input. promote_integers is ignored if dtype is specified.
- out (<u>None</u>) -

Returns: sum along axis – An array with the same shape as a, with the

specified axis removed. If a is a 0-d array, or if axis is None, a

scalar is returned. If an output array is specified, a reference to out

is returned.

Return type: ndarray

swapaxes(axis1, axis2)

Interchange two axes of an array.

LAX-backend implementation of numpy.swapaxes().

The JAX version of this function may in some cases return a copy rather than a view of the input.

Original docstring below.

Parameters: • a (array_like) – Input array.

• axis1 (int) - First axis.

axis2 (<u>int</u>) – Second axis.

Returns: $a_swapped - For NumPy >= 1.10.0$, if a is an ndarray, then a view

of a is returned; otherwise a new array is created. For earlier

NumPy versions a view of a is returned only if the order of the axes

is changed, otherwise the input array is returned.

Return type: ndarray

take(indices, axis=None, out=None, mode=None, unique_indices=False,
indices_are_sorted=False, fill_value=None)

Take elements from an array along an axis.

LAX-backend implementation of numpy.take().

By default, JAX assumes that all indices are in-bounds. Alternative out-of-bound index semantics can be specified via the mode parameter (see below).

Original docstring below.

When axis is not None, this function does the same thing as "fancy" indexing (indexing arrays using arrays); however, it can be easier to use if you need elements along a given axis. A call such as np.take(arr, indices, axis=3) is equivalent to arr[:,:,:,indices,...].

Explained without fancy indexing, this is equivalent to the following use of *ndindex*, which sets each of ii, jj, and kk to a tuple of indices:

```
Ni, Nk = a.shape[:axis], a.shape[axis+1:]
Nj = indices.shape
for ii in ndindex(Ni):
    for jj in ndindex(Nj):
        for kk in ndindex(Nk):
            out[ii + jj + kk] = a[ii + (indices[jj],) + kk]
```

- a (array_like (Ni..., M, Nk...)) The source array.
- **indices** (*array_like* (*Nj...*)) The indices of the values to extract.
- axis (<u>int</u>, optional) The axis over which to select values. By default, the flattened input array is used.
- mode (string, default="fill") Out-of-bounds indexing mode. The
 default mode="fill" returns invalid values (e.g. NaN) for out-of
 bounds indices (see also fill_value below). For more
 discussion of mode options, see jax.numpy.ndarray.at.
- **fill_value** (*optional*) The fill value to return for out-of-bounds slices when mode is 'fill'. Ignored otherwise. Defaults to NaN for inexact types, the largest negative value for signed types, the largest positive value for unsigned types, and True for booleans.
- unique_indices (<u>bool</u>, default=False) If True, the implementation will assume that the indices are unique, which can result in more efficient execution on some backends.
- indices_are_sorted (<u>bool</u>, default=False) If True, the
 implementation will assume that the indices are sorted in
 ascending order, which can lead to more efficient execution on
 some backends.

Returns: out – The returned array has the same type as *a*.

Return type: ndarray (Ni..., Nj..., Nk...)

trace(offset=0, axis1=0, axis2=1, dtype=None, out=None)

Return the sum along diagonals of the array.

LAX-backend implementation of numpy.trace().

Original docstring below.

If a is 2-D, the sum along its diagonal with the given offset is returned, i.e., the sum of elements a[i,i+offset] for all i.

If *a* has more than two dimensions, then the axes specified by axis1 and axis2 are used to determine the 2-D sub-arrays whose traces are returned. The shape of the resulting array is the same as that of *a* with *axis1* and *axis2* removed.

Parameters:

- **a** (array_like) Input array, from which the diagonals are taken.
- offset (<u>int</u>, optional) Offset of the diagonal from the main diagonal. Can be both positive and negative. Defaults to 0.
- axis1 (<u>int</u>, optional) Axes to be used as the first and second axis of the 2-D sub-arrays from which the diagonals should be taken. Defaults are the first two axes of a.
- axis2 (<u>int</u>, optional) Axes to be used as the first and second axis of the 2-D sub-arrays from which the diagonals should be taken. Defaults are the first two axes of a.
- dtype (<u>dtype</u>, optional) Determines the data-type of the
 returned array and of the accumulator where the elements are
 summed. If dtype has the value None and a is of integer type of
 precision less than the default integer precision, then the default
 integer precision is used. Otherwise, the precision is the same as
 that of a.
- out (<u>None</u>) -

Returns: sum_along_diagonals – If a is 2-D, the sum along the diagonal is

returned. If a has larger dimensions, then an array of sums along

diagonals is returned.

Return type: ndarray

unsafe_buffer_pointer()

(self: xla::PyBuffer::pyobject) -> StatusOr[int]

var(axis=None, dtype=None, out=None, ddof=0, keepdims=False, *,
where=None)

Compute the variance along the specified axis.

LAX-backend implementation of numpy.var().

Original docstring below.

Returns the variance of the array elements, a measure of the spread of a distribution. The variance is computed for the flattened array by default, otherwise over the specified axis.

Parameters:

- a (array_like) Array containing numbers whose variance is desired. If a is not an array, a conversion is attempted.
- axis (None or <u>int</u> or tuple of ints, optional) Axis or axes along
 which the variance is computed. The default is to compute the
 variance of the flattened array.
- dtype (data-type, optional) Type to use in computing the variance. For arrays of integer type the default is float64; for arrays of float types it is the same as the array type.
- ddof (<u>int</u>, optional) "Delta Degrees of Freedom": the divisor used in the calculation is N ddof, where N represents the number of elements. By default ddof is zero.
- keepdims (bool, optional) —
 If this is set to True, the axes which are reduced are left in the result as dimensions with size one. With this option, the result will broadcast correctly against the input array.
 If the default value is passed, then keepdims will not be passed through to the var method of sub-classes of ndarray, however any non-default value will be. If the sub-class' method does not implement keepdims any exceptions will be raised.
- where (array_like of bool, optional) Elements to include in the variance. See ~numpy.ufunc.reduce for details.
- out (<u>None</u>) -

Returns: variance – If out=None, returns a new array containing the

variance; otherwise, a reference to the output array is returned.

Return type: ndarray, see dtype parameter above

xla_dynamic_shape()

(self: xla::PyBuffer::pyobject) -> StatusOr[jaxlib.xla_extension.Shape]

xla_shape()

(self: xla::PyBuffer::pyobject) -> jaxlib.xla_extension.Shape

/ The JAX authors

Copyright 2023, The JAX Authors. NumPy and SciPy documentation are copyright the respective authors..