

"""Stax is a small but flexible neural net specification library from scratch.

For an example of its use, see [examples/resnet50.py](#).

"""

```
from __future__ import absolute_import
from __future__ import division
from __future__ import print_function
```

```
import functools
import itertools
import operator as op
```

```
import numpy as onp
import numpy.random as npr
from six.moves import reduce
```

```
from jax import lax
from jax import random
from jax.scipy.misc import logsumexp
import jax.numpy as np
```

```
# Following the convention used in Keras and tf.layers, we use CamelCase for the
# names of layer constructors, like Conv and Relu, while using snake_case for
# other functions, like lax.conv and relu.
```

```
def relu(x): return np.maximum(x, 0.)
def softplus(x): return np.logaddexp(x, 0.)
```

```
def logsoftmax(x, axis=-1):
    """Apply log softmax to an array of logits, log-normalizing along an axis."""
    return x - logsumexp(x, axis, keepdims=True)
```

```
def fastvar(x, axis, keepdims):
    """A fast but less numerically-stable variance calculation than np.var."""
    return np.mean(x**2, axis, keepdims) - np.mean(x, axis, keepdims)**2
```

```
# Initializers
```

```

def randn(stddev=1e-2, rng=npr):
    """An initializer function for random normal coefficients."""
    def init(shape):
        return rng.normal(size=shape, scale=stddev).astype('float32')
    return init

def glorot(out_dim=0, in_dim=1, scale=onp.sqrt(2), rng=npr):
    """An initializer function for random Glorot-scaled coefficients."""
    def init(shape):
        fan_in, fan_out = shape[in_dim], shape[out_dim]
        size = onp.prod(onp.delete(shape, [in_dim, out_dim]))
        std = scale / np.sqrt((fan_in + fan_out) / 2. * size)
        return rng.normal(size=shape, scale=std).astype('float32')
    return init

zeros = functools.partial(np.zeros, dtype='float32')
ones = functools.partial(np.ones, dtype='float32')

```

Layers

Each layer constructor function returns an (init_fun, apply_fun) pair, where

init_fun: takes an input shape and returns an (output_shape, params) pair,

apply_fun: takes params, inputs, and an rng key and applies the layer.

```

def Dense(out_dim, W_init=glorot(), b_init=randn()):
    """Layer constructor function for a dense (fully-connected) layer."""
    def init_fun(input_shape):
        output_shape = input_shape[:-1] + (out_dim,)
        W, b = W_init((input_shape[-1], out_dim)), b_init((out_dim,))
        return output_shape, (W, b)
    def apply_fun(params, inputs, rng=None):
        W, b = params
        return np.dot(inputs, W) + b
    return init_fun, apply_fun

```

```

def GeneralConv(dimension_numbers, out_chan, filter_shape,
                strides=None, padding='VALID', W_init=None, b_init=randn(1e-6)):

```

```

"""Layer construction function for a general convolution layer."""
lhs_spec, rhs_spec, out_spec = dimension_numbers
one = (1,) * len(filter_shape)
strides = strides or one
W_init = W_init or glorot(rhs_spec.index('O'), rhs_spec.index('I'))
def init_fun(input_shape):
    filter_shape_iter = iter(filter_shape)
    kernel_shape = [out_chan if c == 'O' else
                     input_shape[lhs_spec.index('C')] if c == 'I' else
                     next(filter_shape_iter) for c in rhs_spec]
    output_shape = lax.conv_general_shape_tuple(
        input_shape, kernel_shape, strides, padding, dimension_numbers)
    bias_shape = [out_chan if c == 'C' else 1 for c in out_spec]
    bias_shape = tuple(itertools.dropwhile(lambda x: x == 1, bias_shape))
    W, b = W_init(kernel_shape), b_init(bias_shape)
    return output_shape, (W, b)
def apply_fun(params, inputs, rng=None):
    W, b = params
    return lax.conv_general_dilated(inputs, W, strides, padding, one, one,
                                    dimension_numbers) + b

return init_fun, apply_fun
Conv = functools.partial(GeneralConv, ('NHWC', 'HWIO', 'NHWC'))

def BatchNorm(axis=(0, 1, 2), epsilon=1e-5, center=True, scale=True,
              beta_init=zeros, gamma_init=ones):
    """Layer construction function for a batch normalization layer."""
    _beta_init = lambda shape: beta_init(shape) if center else ()
    _gamma_init = lambda shape: gamma_init(shape) if scale else ()
    axis = (axis,) if np.isscalar(axis) else axis
    def init_fun(input_shape):
        shape = (1 if i in axis else d for i, d in enumerate(input_shape))
        shape = tuple(itertools.dropwhile(lambda x: x == 1, shape))
        beta, gamma = _beta_init(shape), _gamma_init(shape)
        return input_shape, (beta, gamma)
    def apply_fun(params, x, rng=None):
        beta, gamma = params
        mean, var = np.mean(x, axis, keepdims=True), fastvar(x, axis, keepdims=True)
        z = (x - mean) / (var + epsilon)**2
        if center and scale: return gamma * z + beta
        if center: return z + beta

```

```

    if scale: return gamma * z
    return z
return init_fun, apply_fun

def _elemwise_no_params(fun, **kwargs):
    init_fun = lambda input_shape: (input_shape, ())
    apply_fun = lambda params, inputs, rng=None: fun(inputs, **kwargs)
    return init_fun, apply_fun

Tanh = _elemwise_no_params(np.tanh)
Relu = _elemwise_no_params(relu)
LogSoftmax = _elemwise_no_params(logsoftmax, axis=-1)
Softplus = _elemwise_no_params(softplus)

def _pooling_layer(reducer, init_val, rescaler=None):
    def PoolingLayer(window_shape, strides=None, padding='VALID'):
        """Layer construction function for a pooling layer."""
        strides = strides or (1,) * len(window_shape)
        rescale = rescaler(window_shape, strides, padding) if rescaler else None
        dims = (1,) + window_shape + (1,)  # NHWC
        strides = (1,) + strides + (1,)
        def init_fun(input_shape):
            out_shape = lax.reduce_window_shape_tuple(input_shape, dims, strides, padding)
            return out_shape, ()
        def apply_fun(params, inputs, rng=None):
            out = lax.reduce_window(inputs, init_val, reducer, dims, strides, padding)
            return rescale(out, inputs) if rescale else out
        return init_fun, apply_fun
    return PoolingLayer

MaxPool = _pooling_layer(lax.max, -np.inf)
SumPool = _pooling_layer(lax.add, 0.)

def _normalize_by_window_size(dims, strides, padding):
    def rescale(outputs, inputs):
        one = np.ones(inputs.shape[1:3], dtype=inputs.dtype)
        window_sizes = lax.reduce_window(one, 0., lax.add, dims, strides, padding)
        return outputs / window_sizes
    return rescale

AvgPool = _pooling_layer(lax.add, 0., _normalize_by_window_size)

```

```
def Flatten():
    """Layer construction function for flattening all but the leading dim."""
    def init_fun(input_shape):
        output_shape = input_shape[0], reduce(op.mul, input_shape[1:], 1)
        return output_shape, ()
    def apply_fun(params, inputs, rng=None):
        return np.reshape(inputs, (inputs.shape[0], -1))
    return init_fun, apply_fun
Flatten = Flatten()
```

```
def Identity():
    """Layer construction function for an identity layer."""
    init_fun = lambda input_shape: (input_shape, ())
    apply_fun = lambda params, inputs, rng=None: inputs
    return init_fun, apply_fun
Identity = Identity()
```

```
def FanOut(num):
    """Layer construction function for a fan-out layer."""
    init_fun = lambda input_shape: ([input_shape] * num, ())
    apply_fun = lambda params, inputs, rng=None: [inputs] * num
    return init_fun, apply_fun
```

```
def FanInSum():
    """Layer construction function for a fan-in sum layer."""
    init_fun = lambda input_shape: (input_shape[0], ())
    apply_fun = lambda params, inputs, rng=None: sum(inputs)
    return init_fun, apply_fun
FanInSum = FanInSum()
```

```
def Dropout(rate, mode='train'):
    """Layer construction function for a dropout layer with given rate."""
    def init_fun(input_shape):
        return input_shape, ()
    def apply_fun(params, inputs, rng):
```

```

if mode == 'train':
    keep = random.bernoulli(rng, rate, inputs.shape)
    return np.where(keep, inputs / rate, 0)
else:
    return inputs
return init_fun, apply_fun

```

Composing layers via combinators

```
def serial(*layers):
```

```
    """Combinator for composing layers in serial.
```

Args:

*layers: a sequence of layers, each an (init_fun, apply_fun) pair.

Returns:

A new layer, meaning an (init_fun, apply_fun) pair, representing the serial composition of the given sequence of layers.

```
    """
```

```
    nlayers = len(layers)
```

```
    init_funs, apply_funs = zip(*layers)
```

```
    def init_fun(input_shape):
```

```
        params = []
```

```
        for init_fun in init_funs:
```

```
            input_shape, param = init_fun(input_shape)
```

```
            params.append(param)
```

```
        return input_shape, params
```

```
    def apply_fun(params, inputs, rng=None):
```

```
        rngs = random.split(rng, nlayers) if rng is not None else (None,) * nlayers
```

```
        for fun, param, rng in zip(apply_funs, params, rngs):
```

```
            inputs = fun(param, inputs, rng)
```

```
        return inputs
```

```
    return init_fun, apply_fun
```

```
def parallel(*layers):
```

```
    """Combinator for composing layers in parallel.
```

The layer resulting from this combinator is often used with the FanOut and

FanInSum layers.

Args:

*layers: a sequence of layers, each an (init_fun, apply_fun) pair.

Returns:

A new layer, meaning an (init_fun, apply_fun) pair, representing the parallel composition of the given sequence of layers. In particular, the returned layer takes a sequence of inputs and returns a sequence of outputs with the same length as the argument `layers`.

```
"""
nlayers = len(layers)
init_funs, apply_funs = zip(*layers)
def init_fun(input_shape):
    return zip(*[init(shape) for init, shape in zip(init_funs, input_shape)])
def apply_fun(params, inputs, rng=None):
    rngs = random.split(rng, nlayers) if rng is not None else (None,) * nlayers
    return [f(p, x, r) for f, p, x, r in zip(apply_funs, params, inputs, rngs)]
return init_fun, apply_fun
```

```
def shape_dependent(make_layer):
```

```
    """Combinator to delay layer constructor pair until input shapes are known.
```

Args:

make_layer: a one-argument function that takes an input shape as an argument (a tuple of positive integers) and returns an (init_fun, apply_fun) pair.

Returns:

A new layer, meaning an (init_fun, apply_fun) pair, representing the same layer as returned by `make_layer` but with its construction delayed until input shapes are known.

```
"""
def init_fun(input_shape):
    return make_layer(input_shape)[0](input_shape)
def apply_fun(params, inputs, rng=None):
    return make_layer(inputs.shape)[1](params, inputs, rng)
return init_fun, apply_fun
```